

CONFIDENTIAL COMPUTING

Some basic concepts and terminology

Prof. Dr. Christof Fetzer

Confidential Computing: Application Domains

CLOUD COMPUTING

- Protect cloud-native applications
- Protect AI applications
- Protect cloud services itself

EDGE CLOUD

- Protect applications and cloud
 - despite limited physical security

AIR-GAPPED SYSTEMS

- Protect applications despite having no connectivity

EMBEDDED SYSTEMS

- Protect devices with limited connectivity & limited physical security

eHEALTH DOMAIN

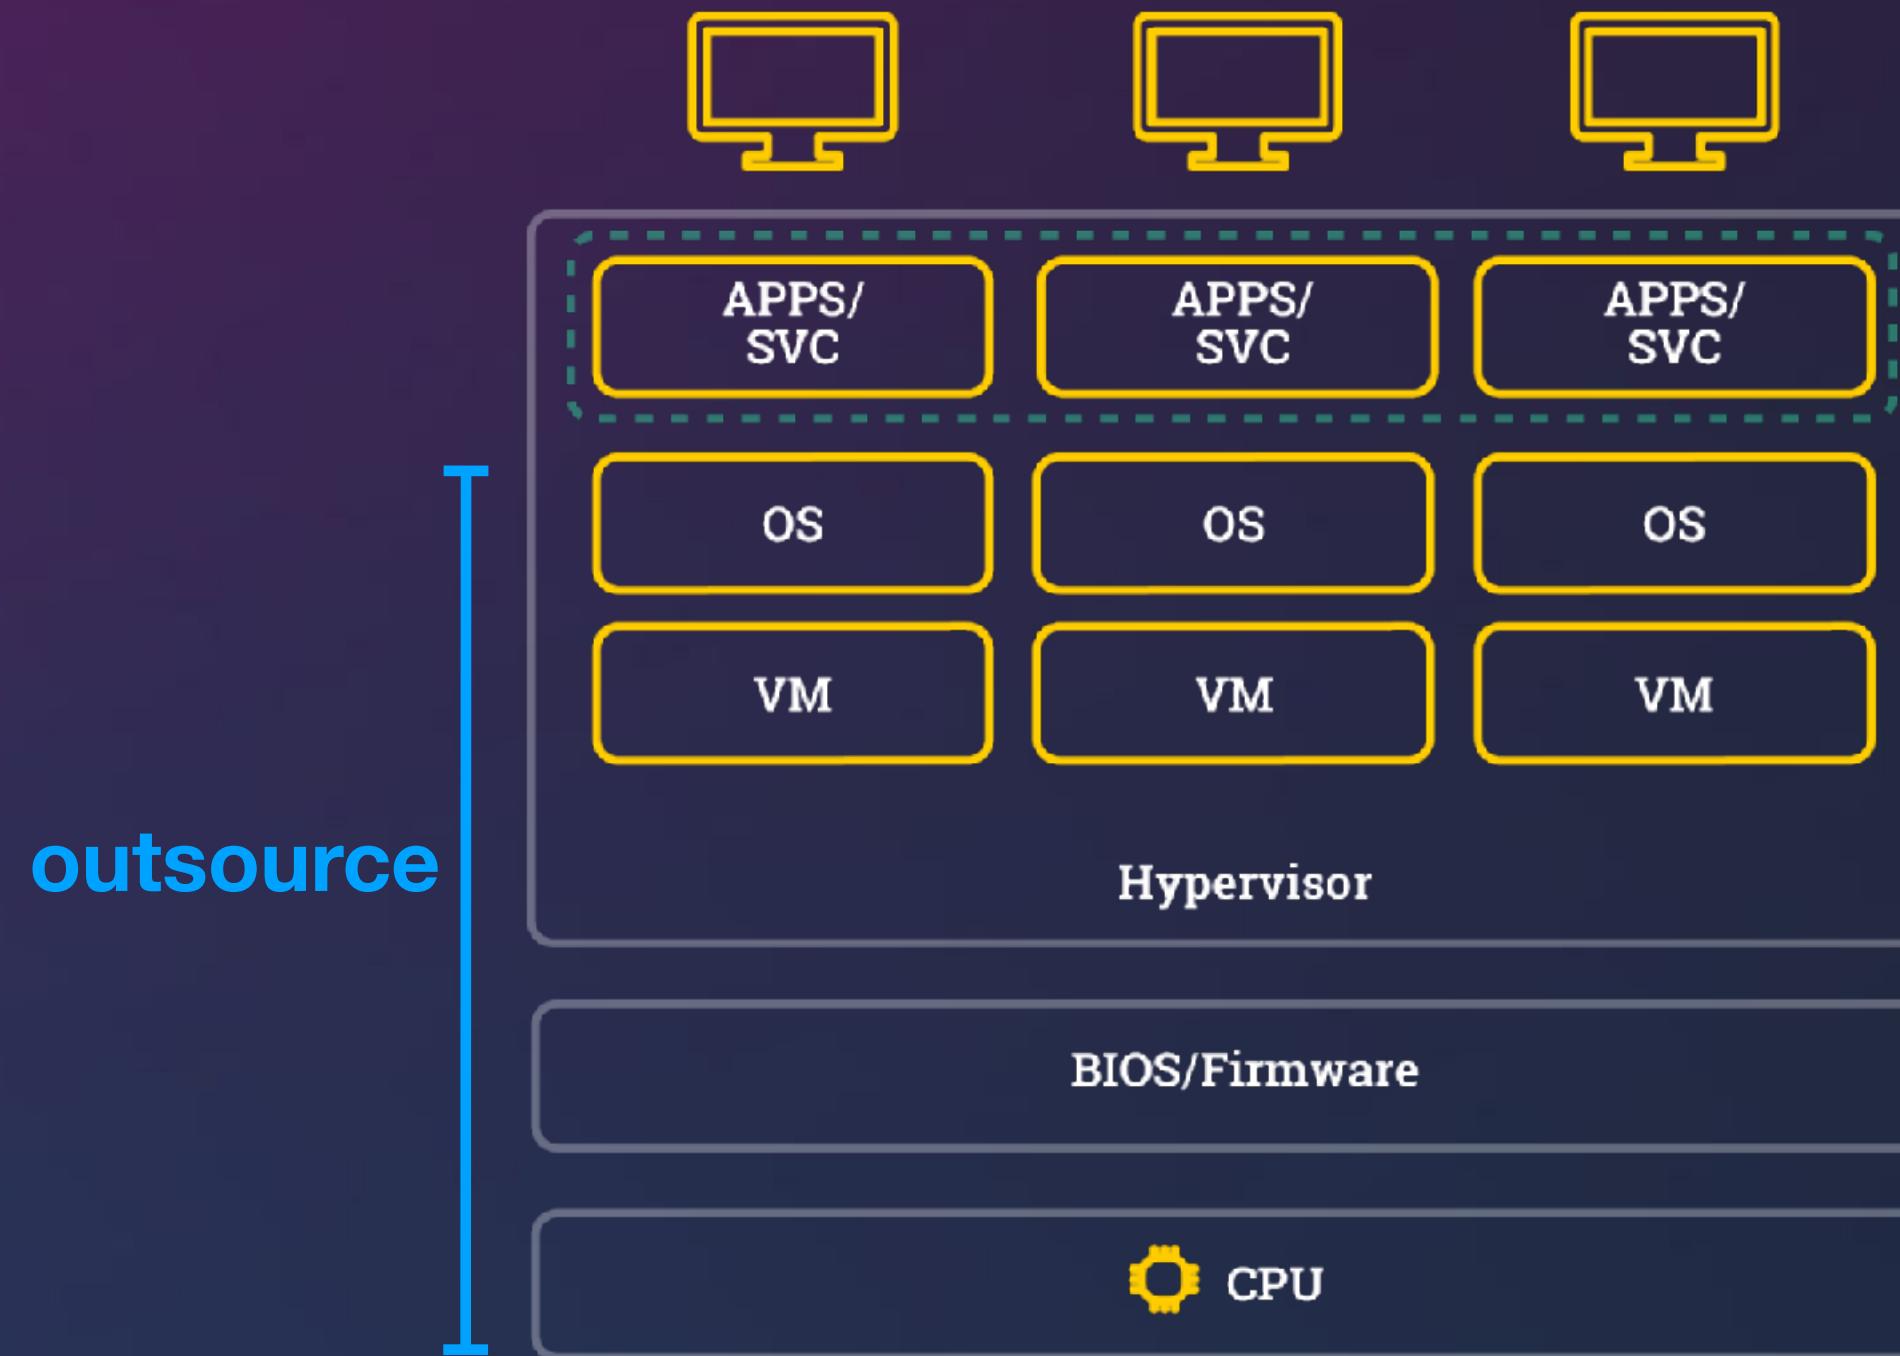
- e.g., protect patient data

Cloud: Outsource Infrastructure

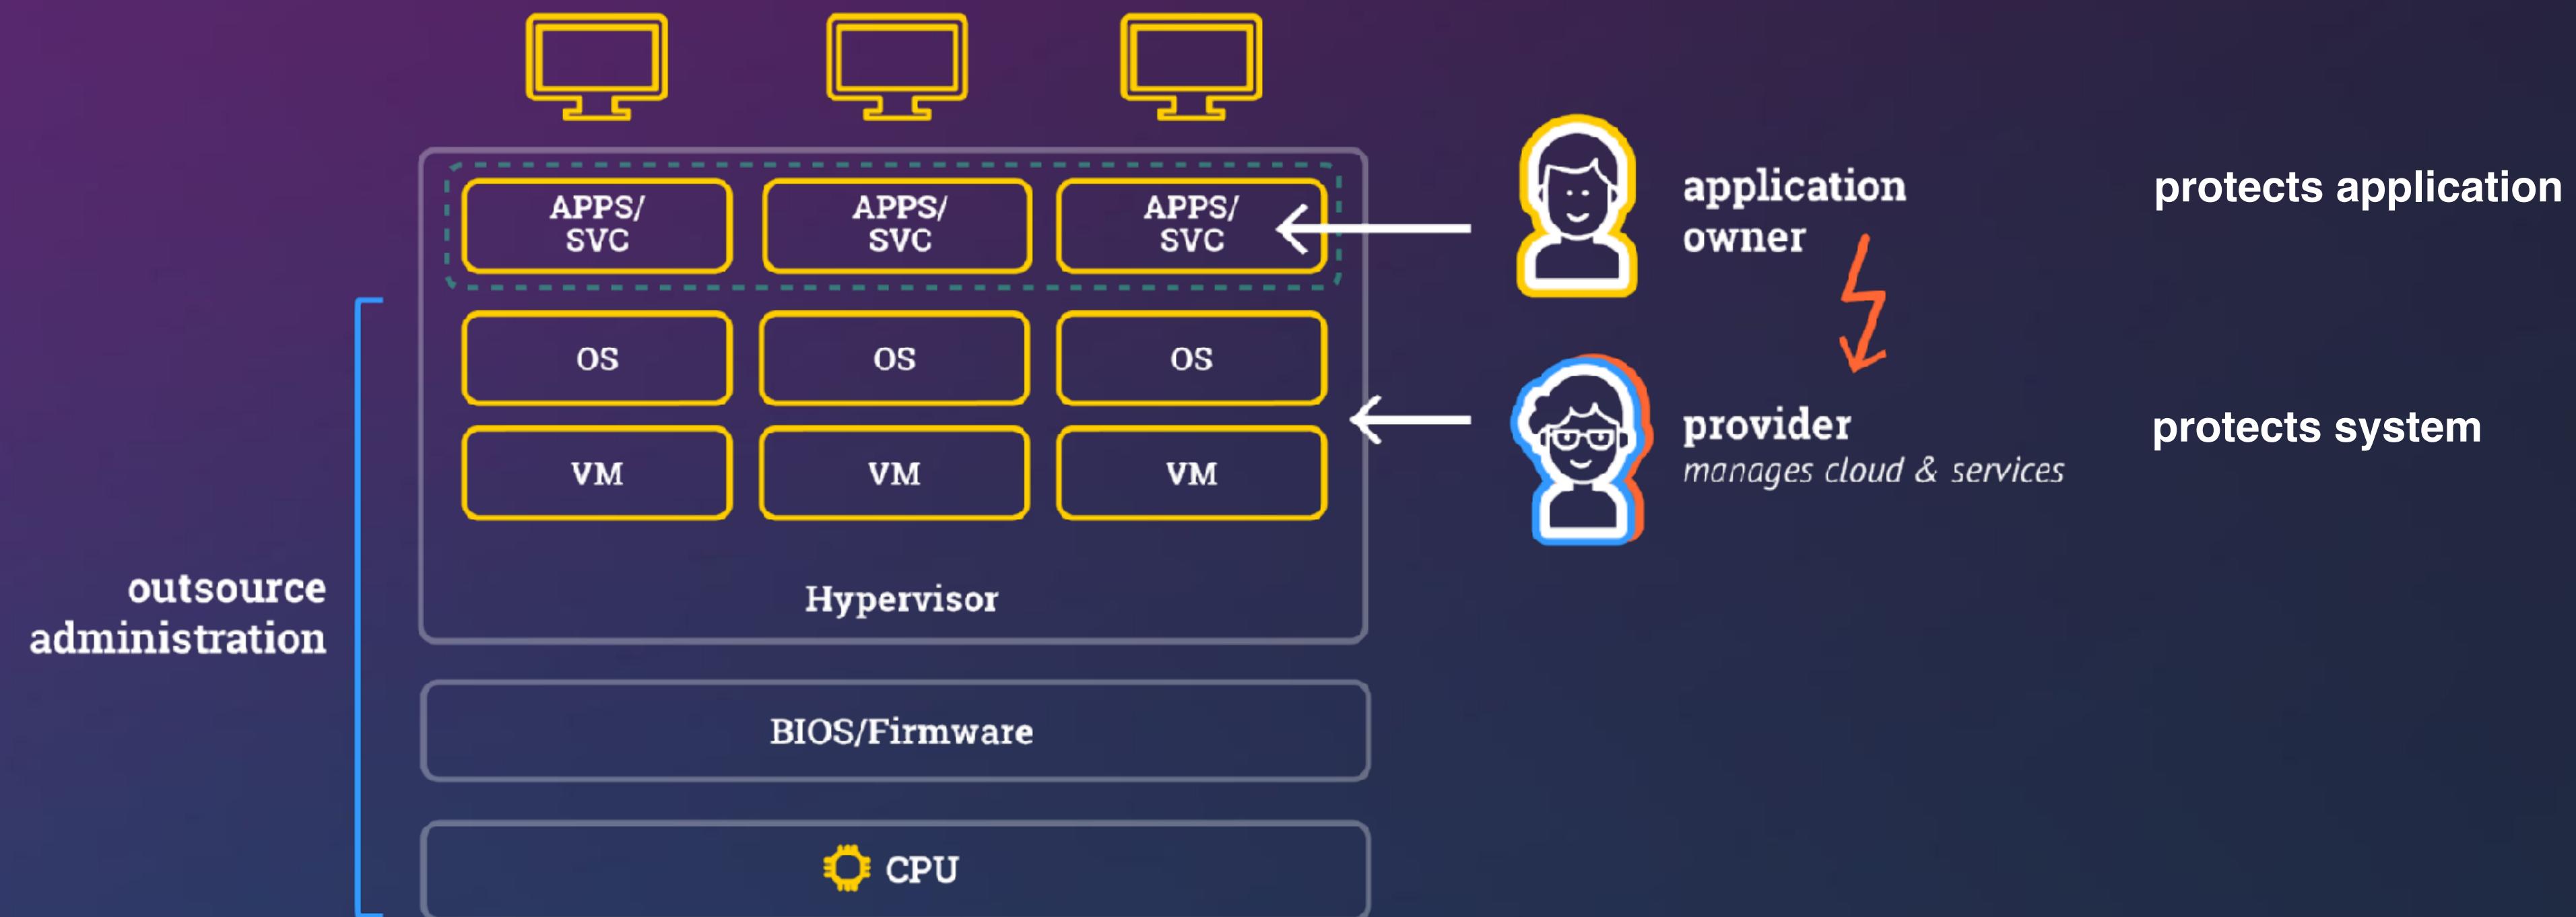


Operating one's own computing infrastructure is not easy:

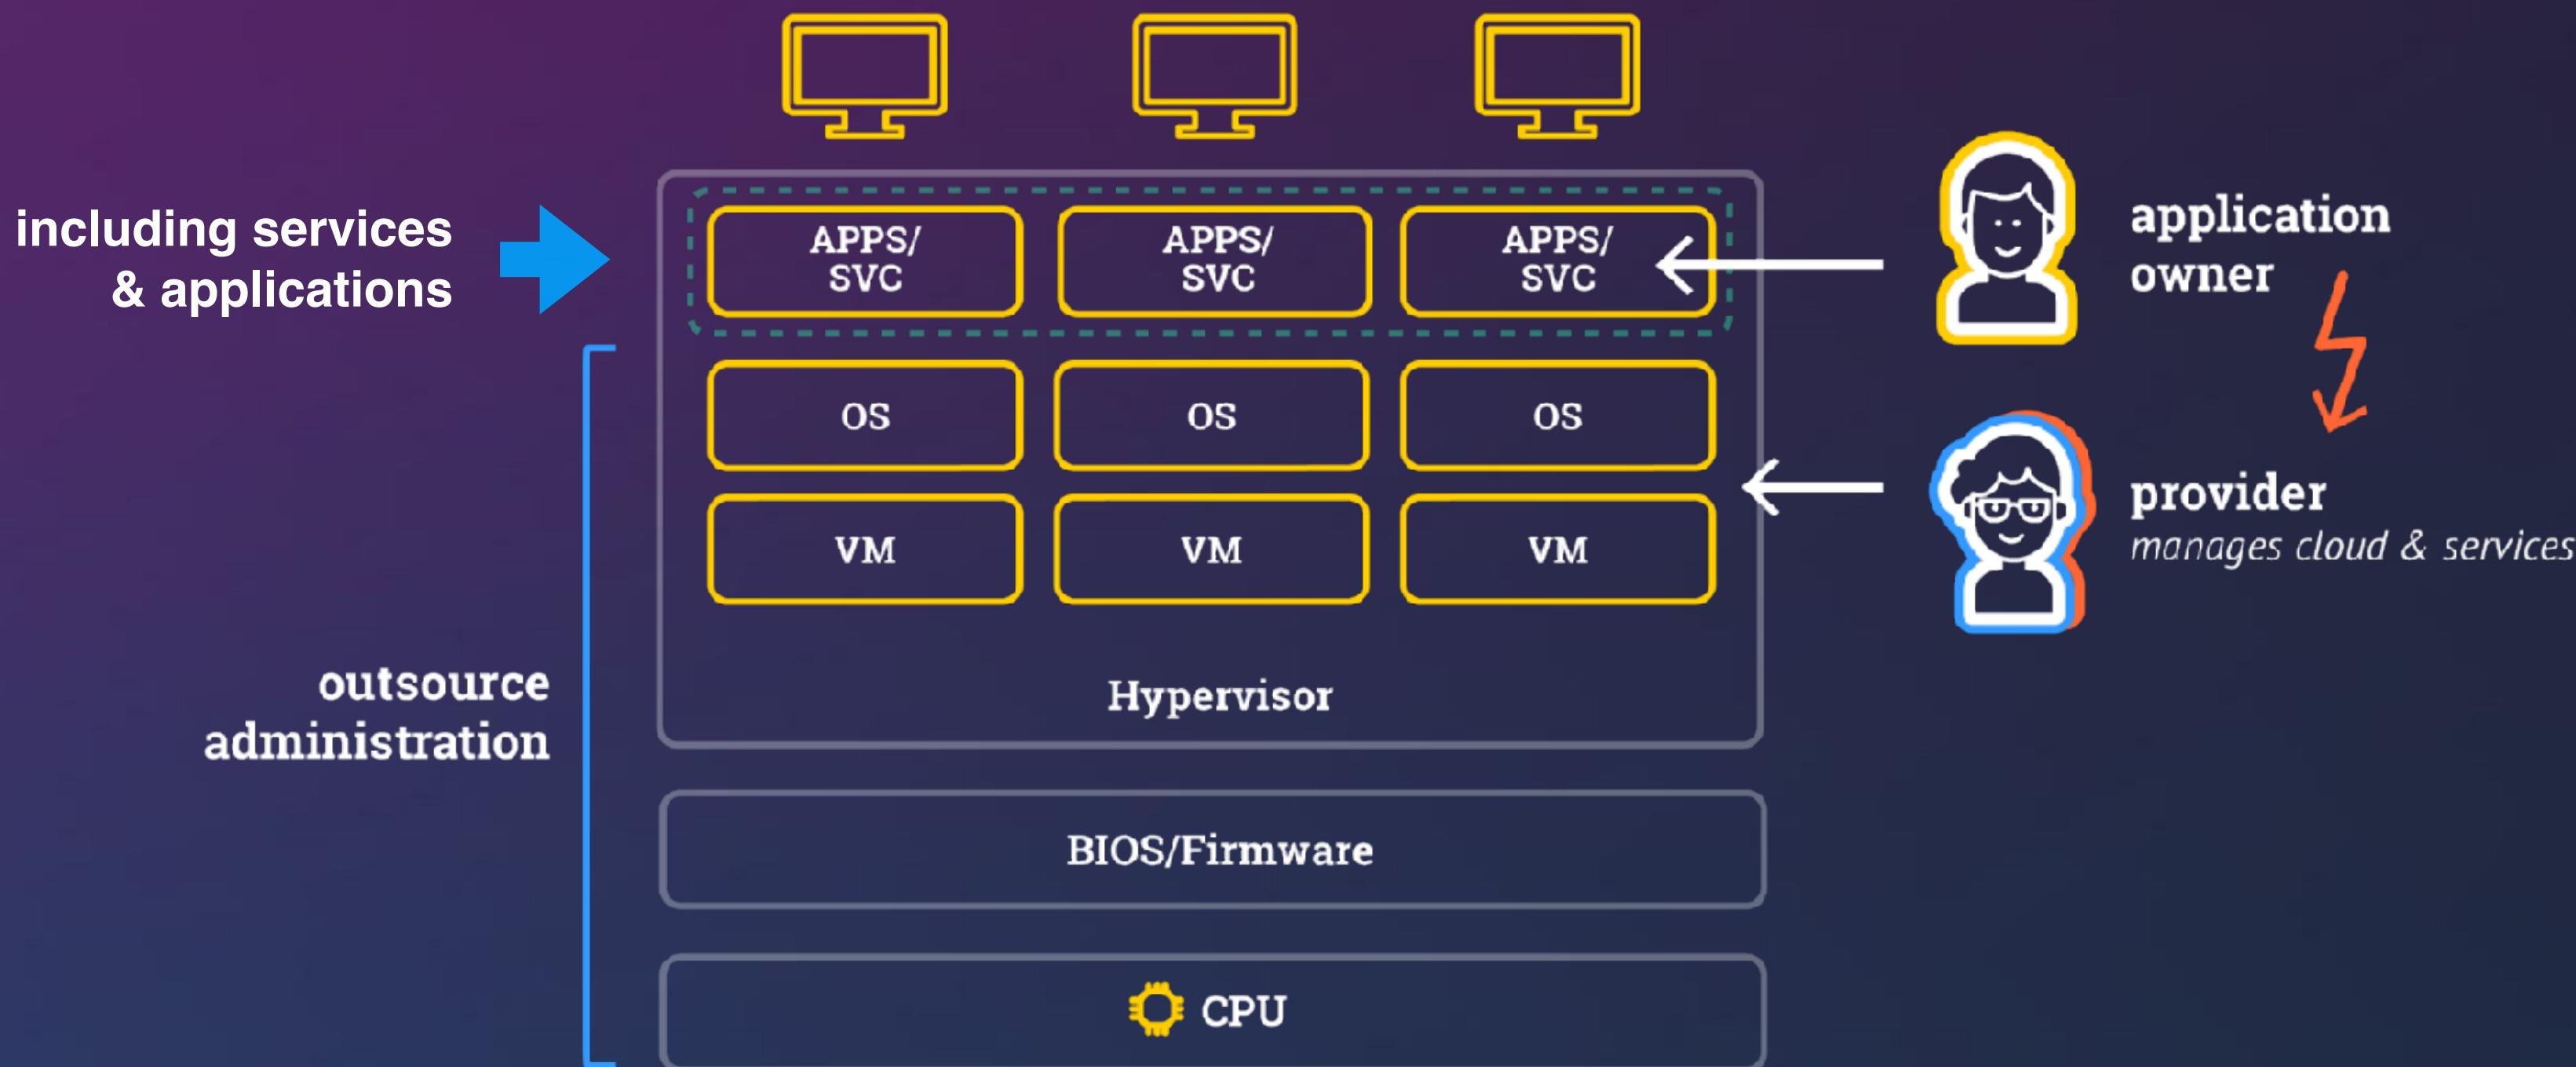
- data center
- hypervisors
- operating systems
- Kubernetes
- services
- ...



Different Stakeholders



Confidential Managed Services

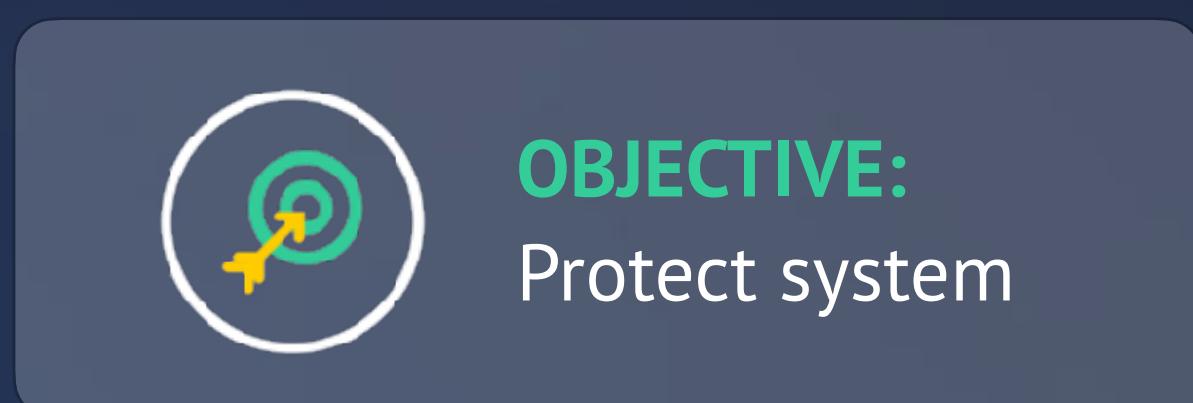
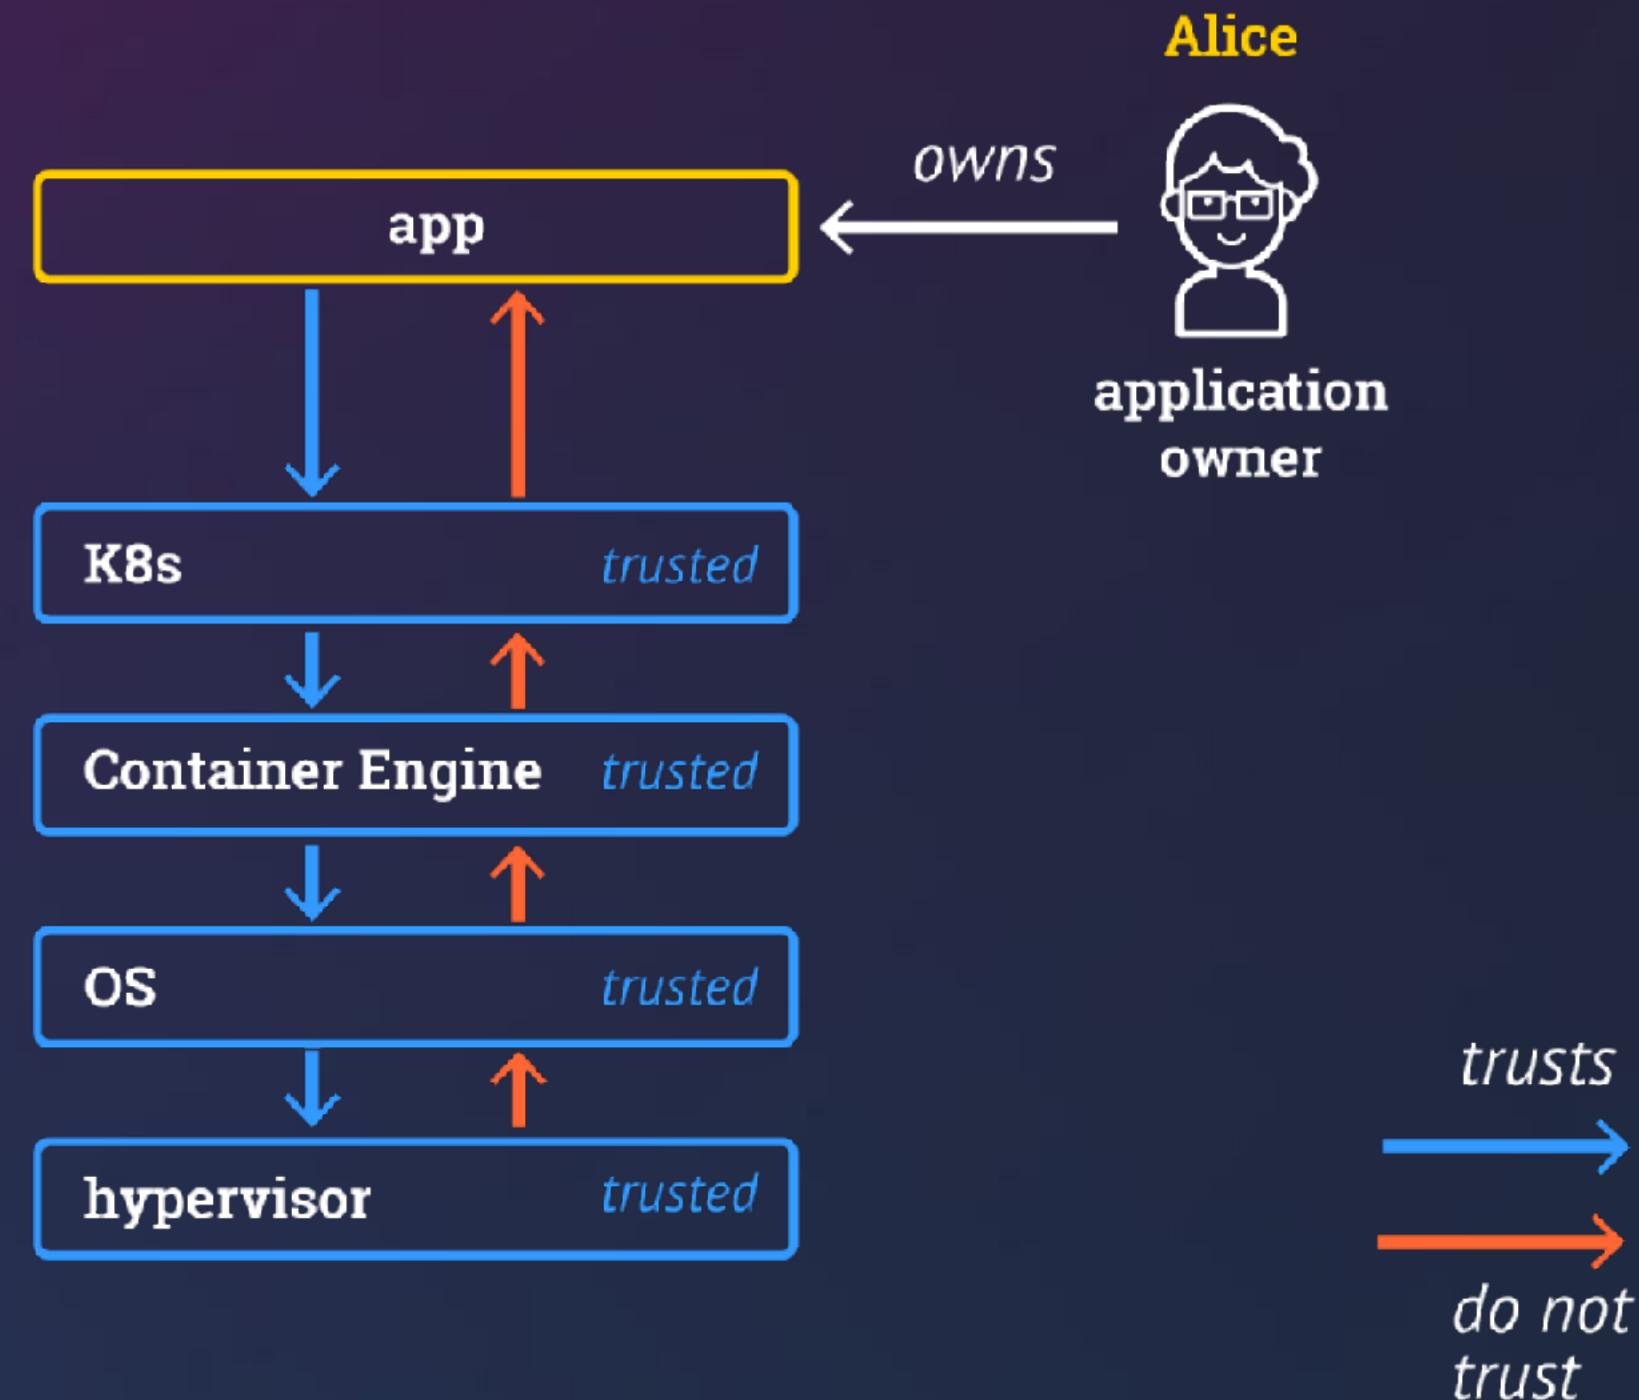


OBJECTIVE:

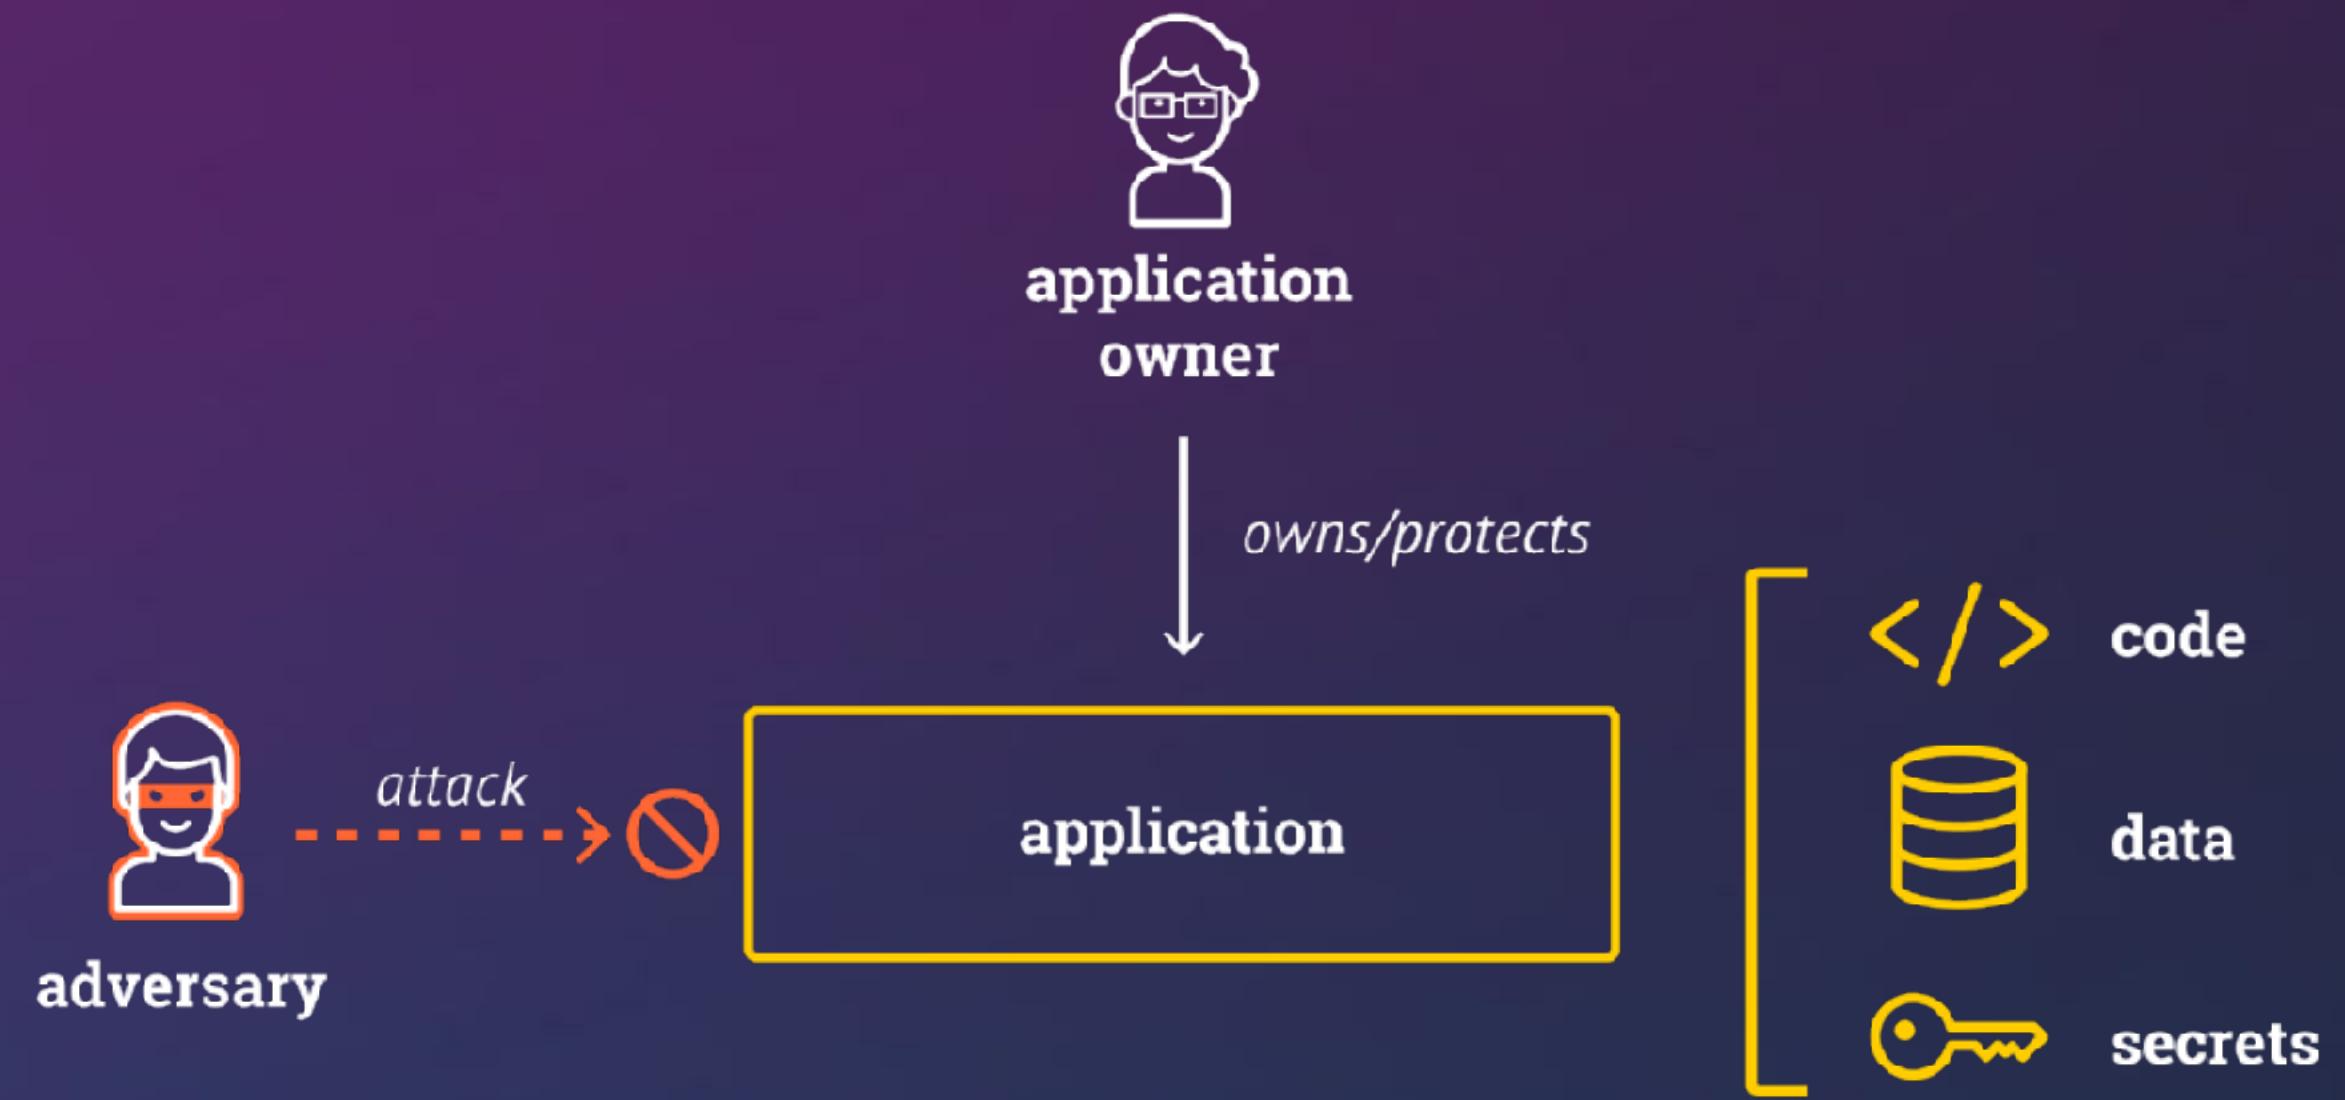
Outsource management of services and applications
provider must not have access to data / code / secrets.

Systems Security

- Systems are structured in layers
 - like operating system and hypervisor
- Typically, systems security is bottom up
 - layer i does not trust layer $i+1$
 - but layer $i+1$ trusts layer i
- Examples:
 - the operating system trusts the hypervisor
 - but the hypervisor does not trust the operating system

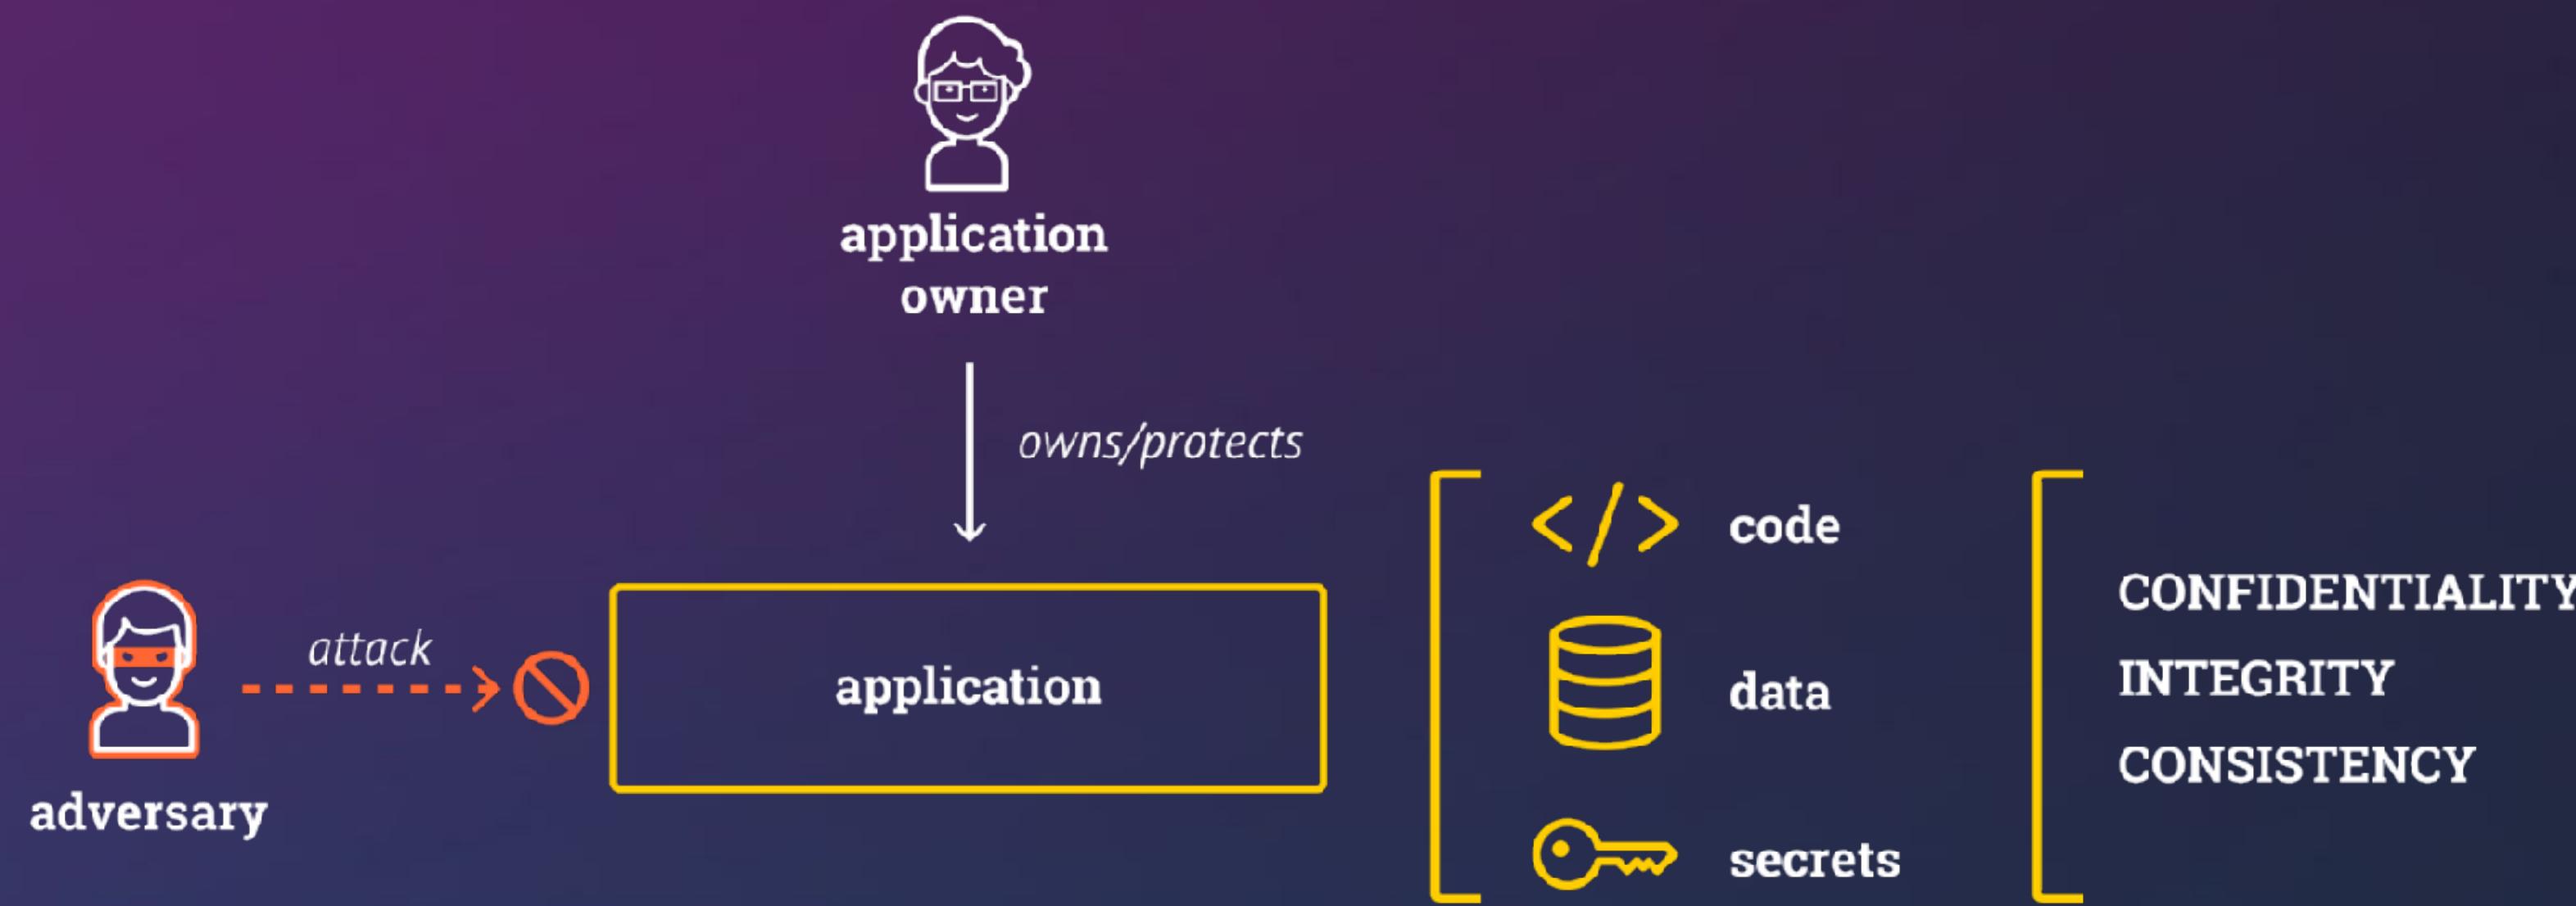


Confidential Computing: Protect Applications



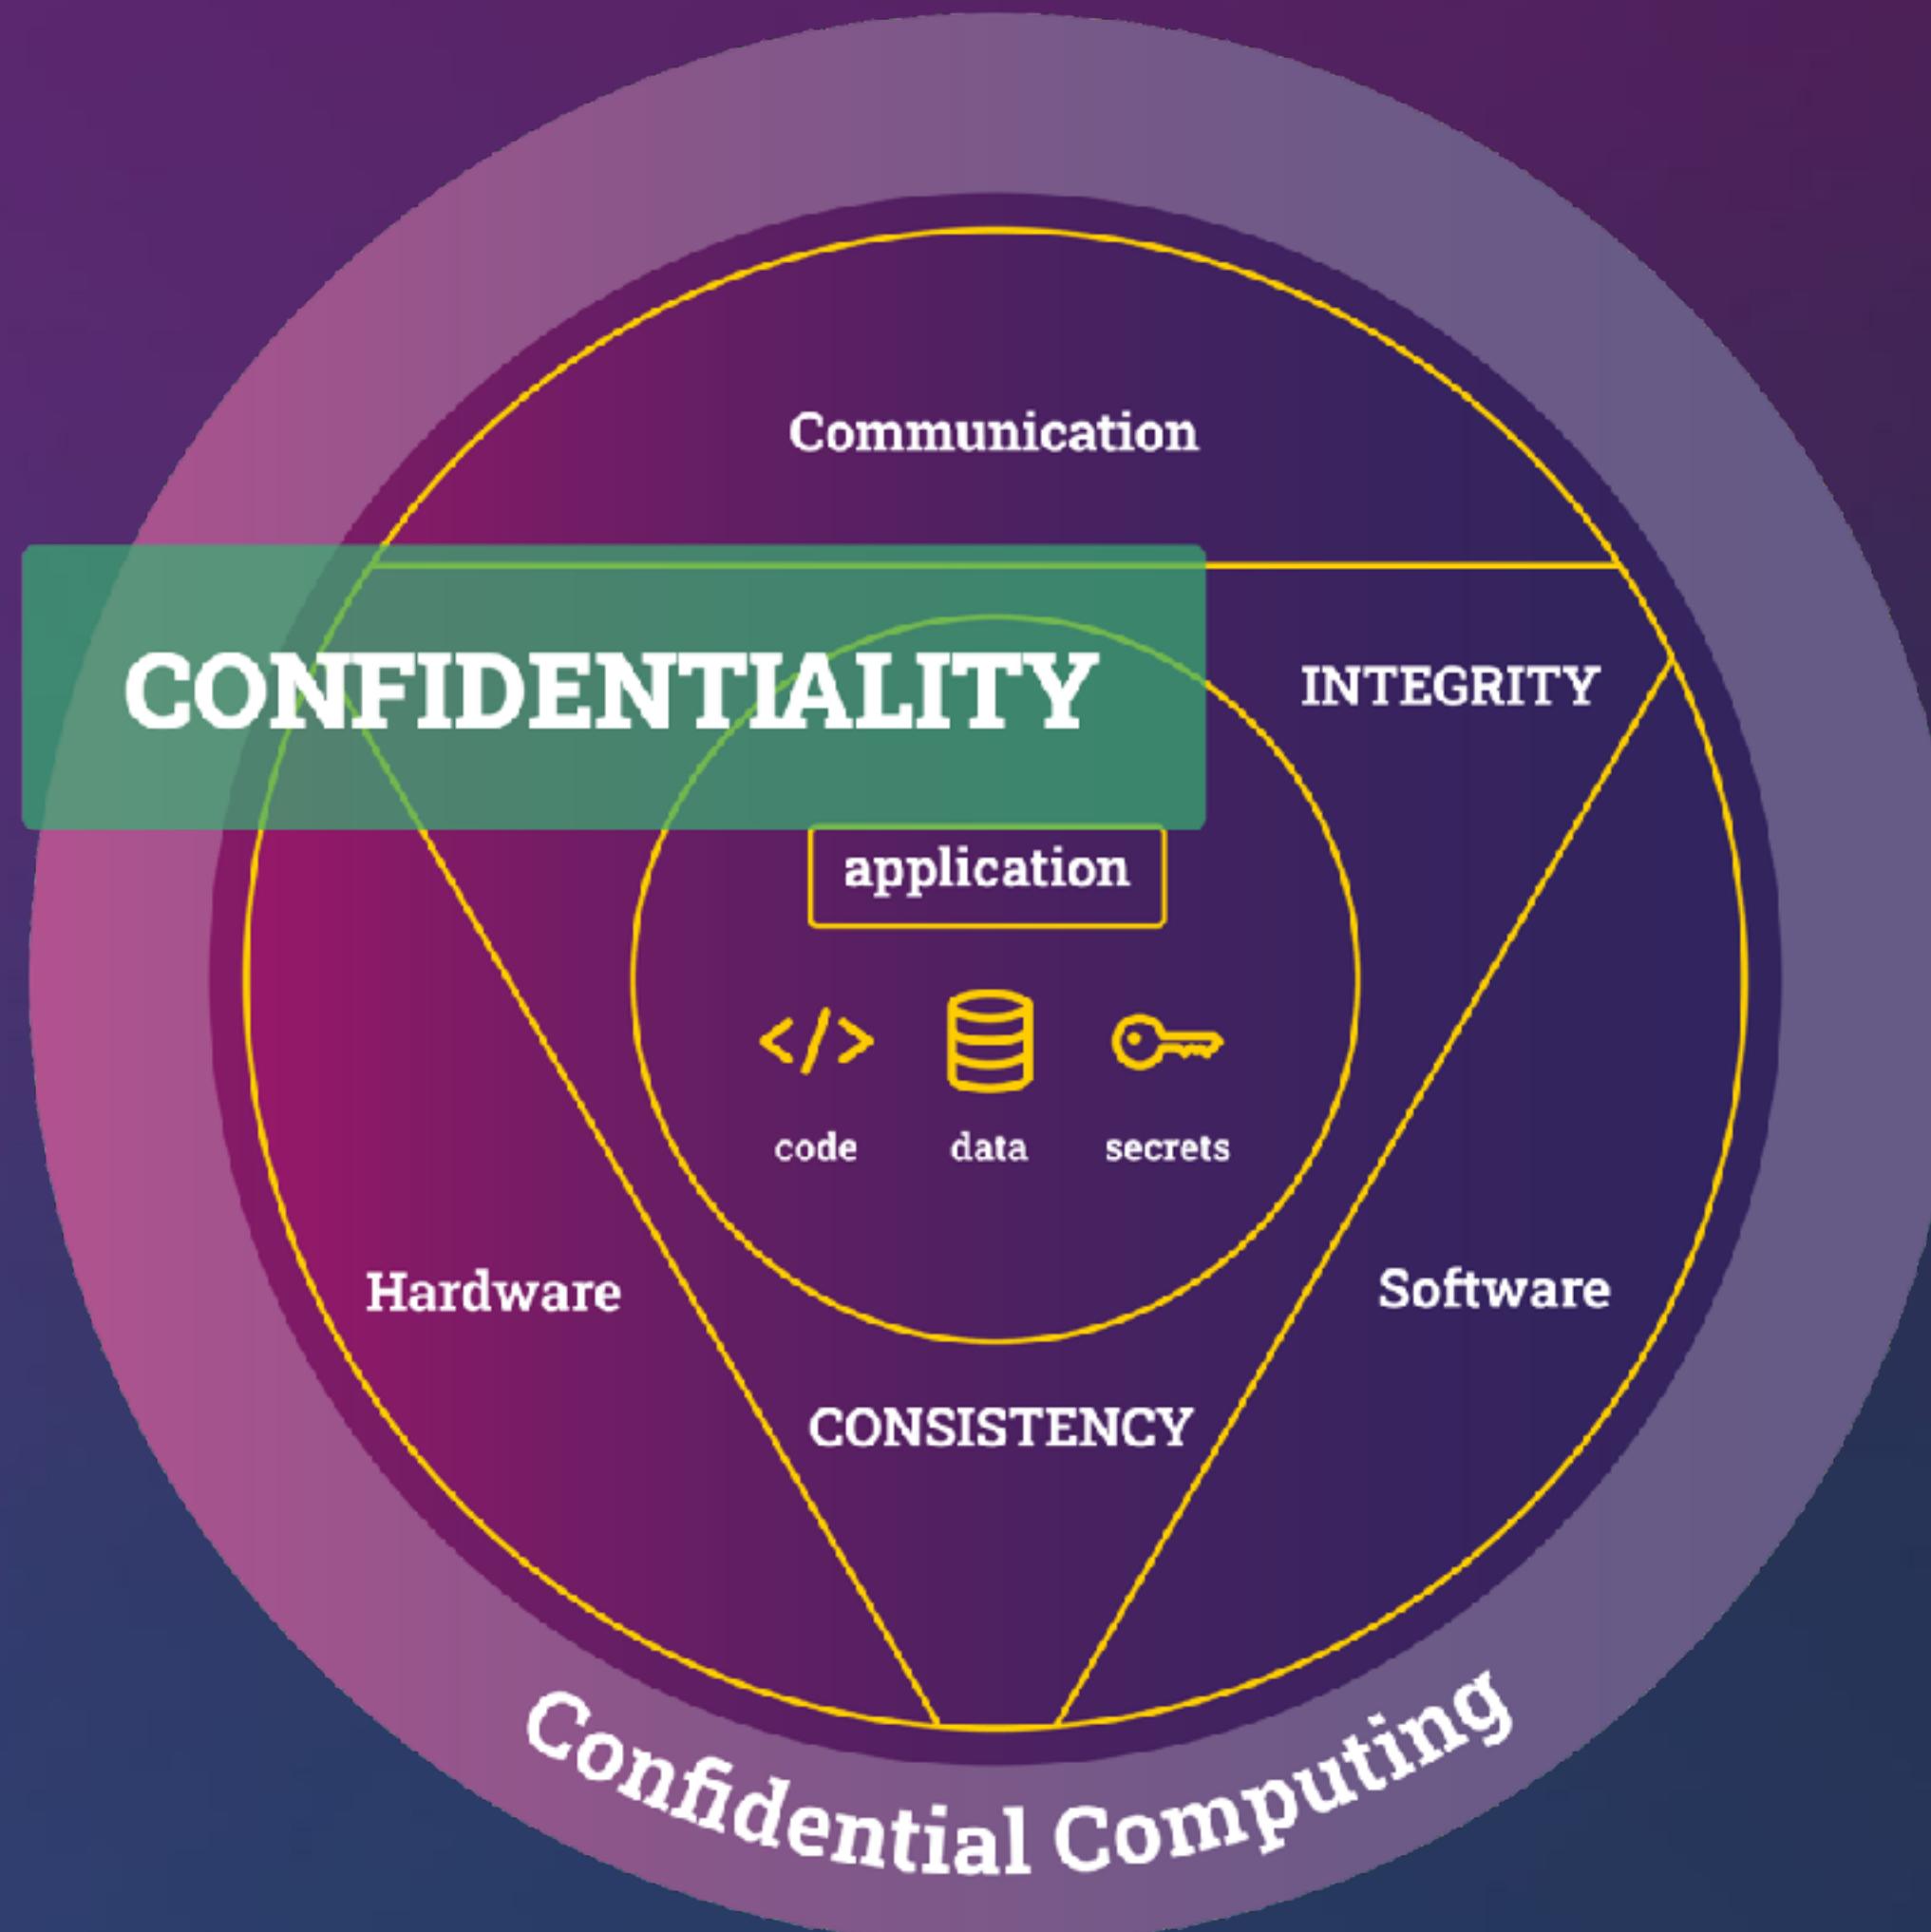
„application-oriented security“

Protect Code, Data, and Secrets



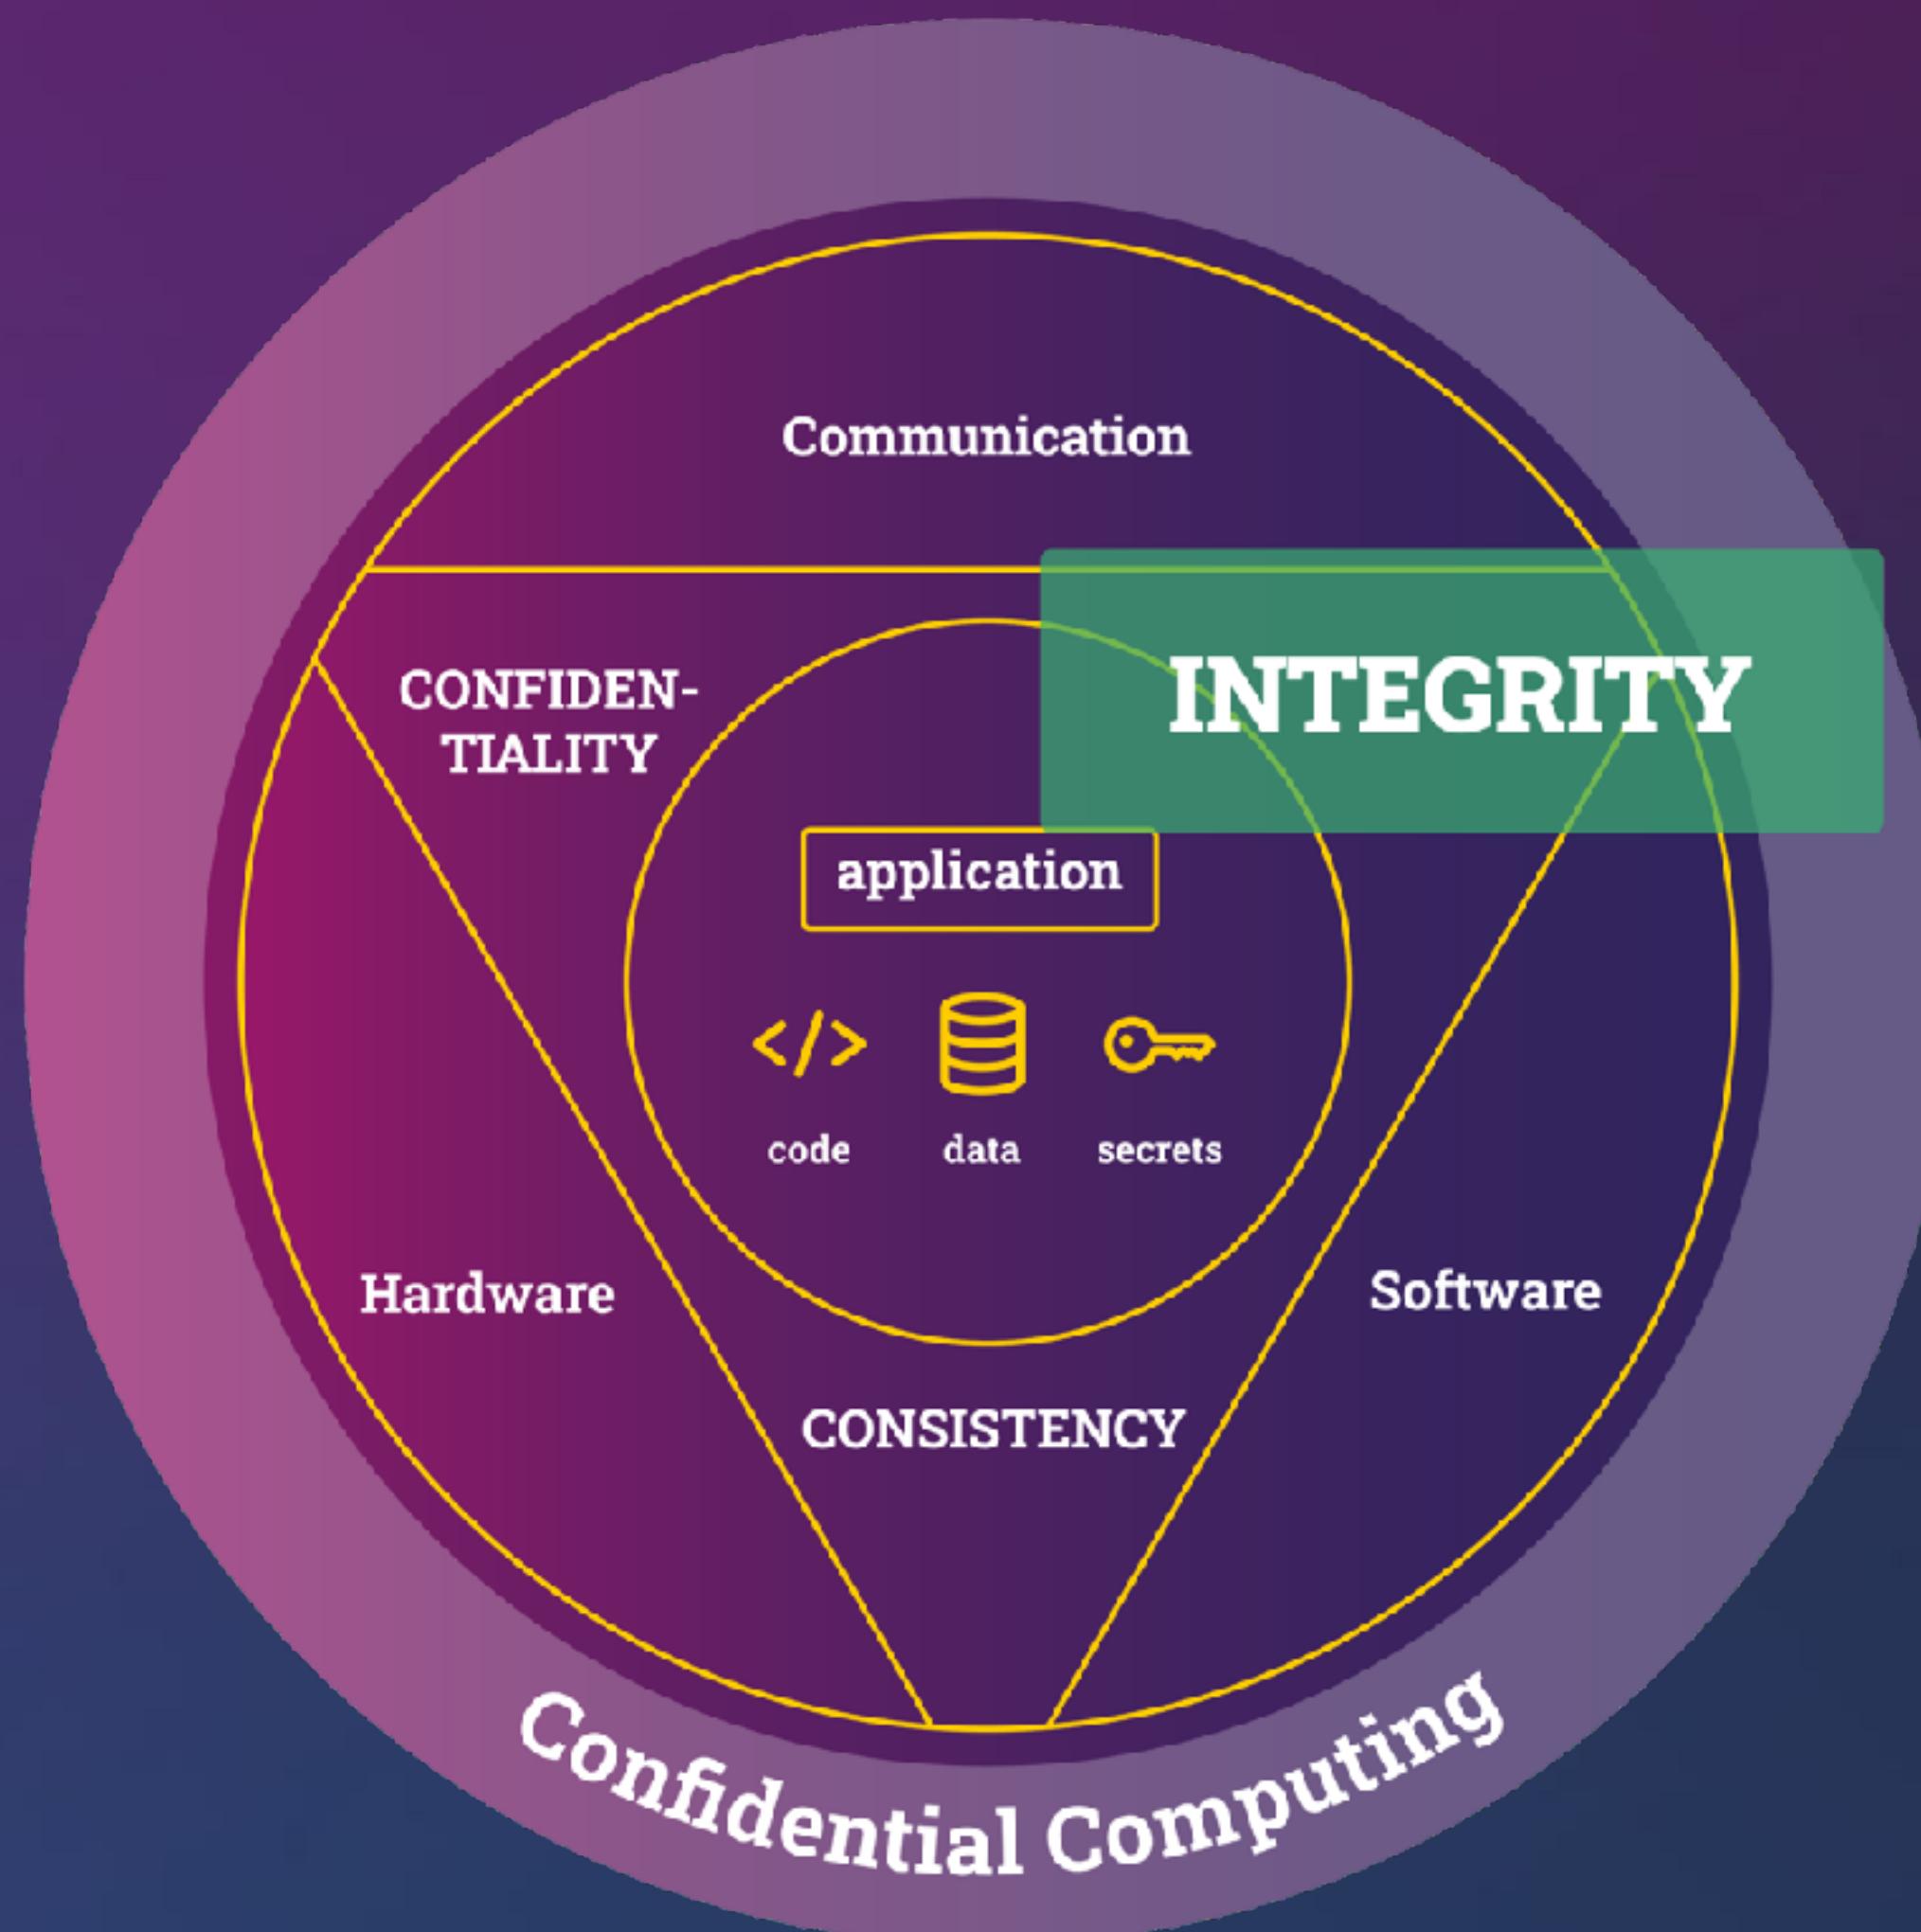
„application-oriented security“

Confidentiality



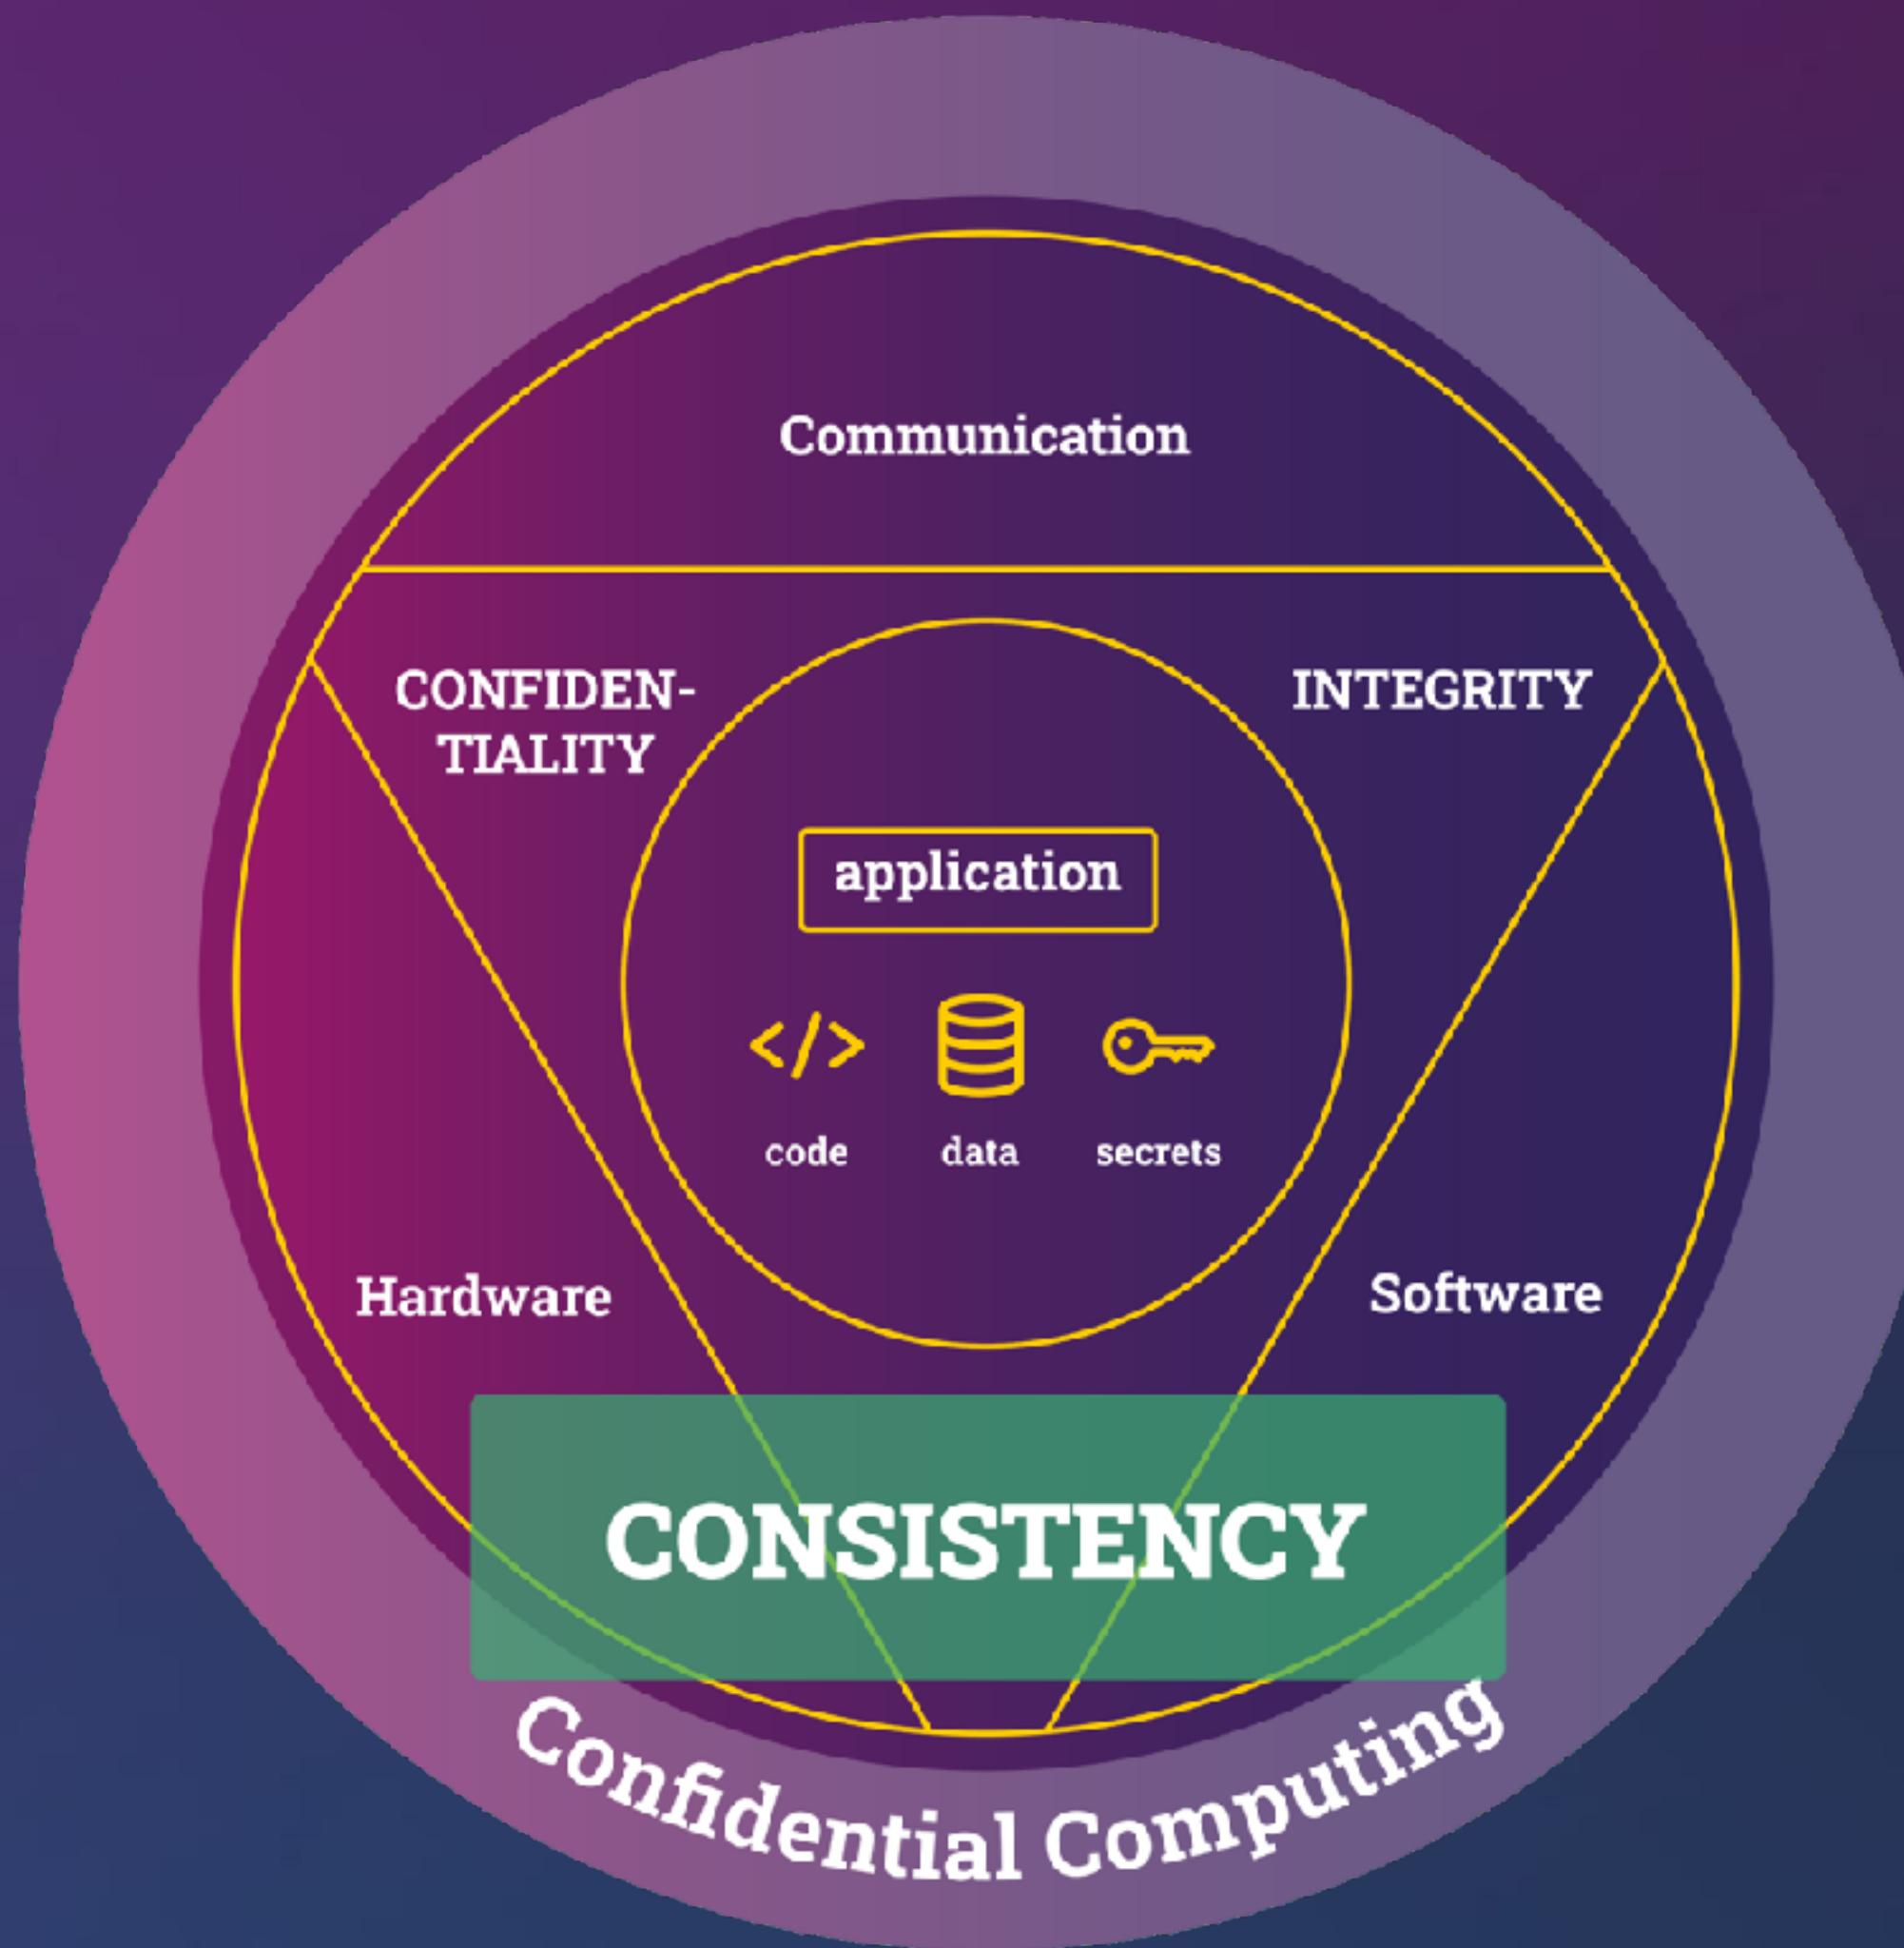
Confidentiality is the property, that information (data, code, secrets) is not made available or disclosed to **unauthorized** individuals, entities, or processes.

Integrity



Integrity means that information (i.e., data, code, secrets) cannot be modified in an **unauthorized** or **undetected** manner.

Consistency



Consistency means that one always reads the latest information (i.e., data, code, secrets) written by an **authorized** entity.

Detect if an adversary would provide an old copies (which are correctly encrypted but that have been updated).

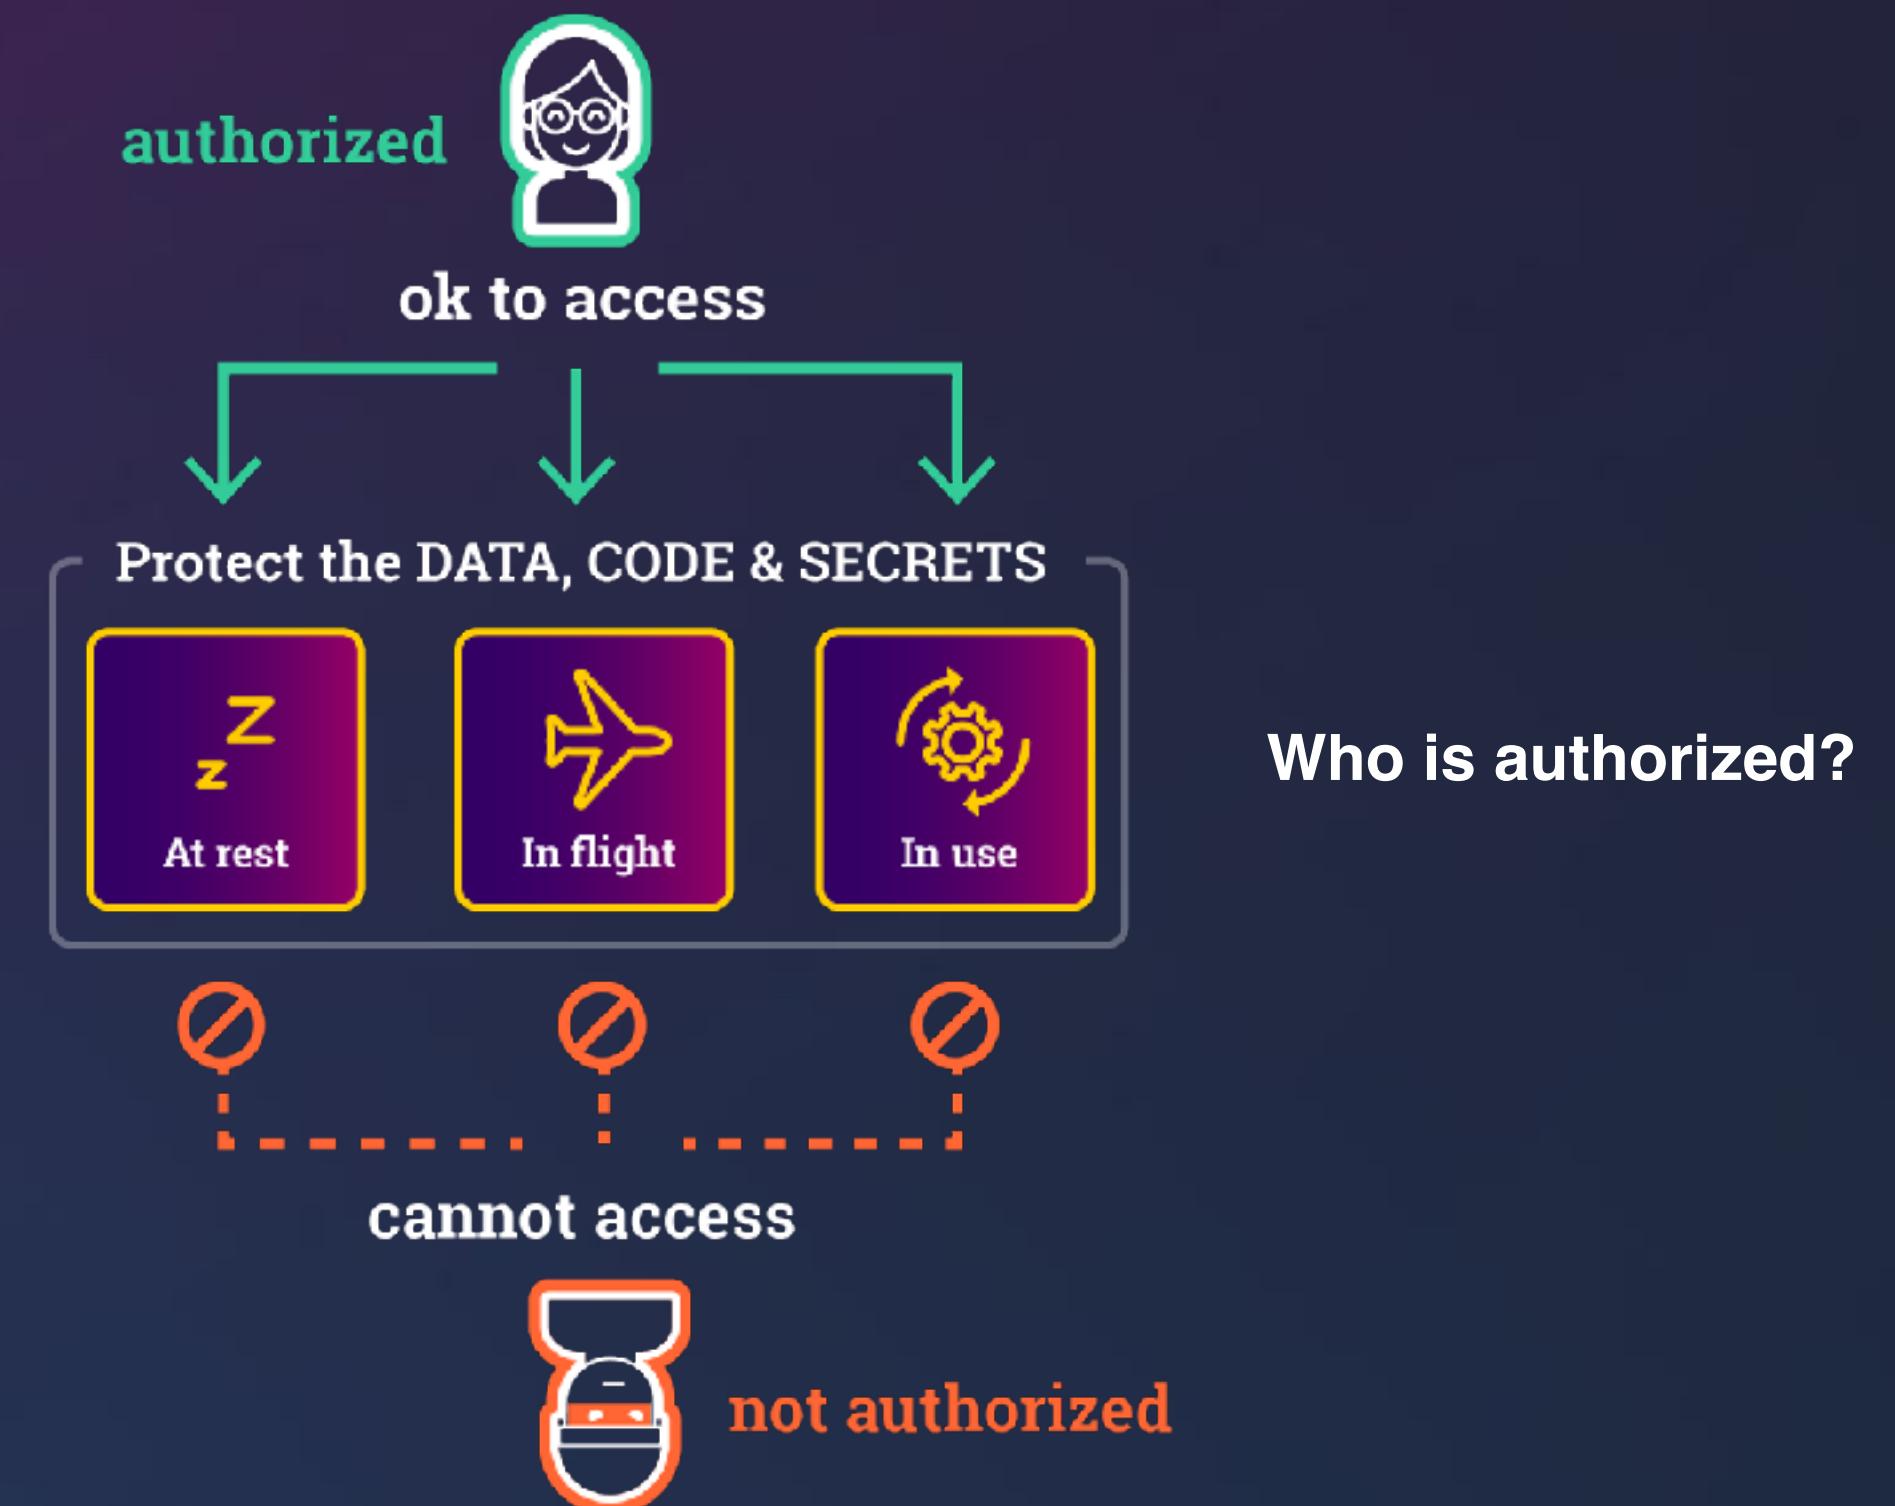
PROBLEM:



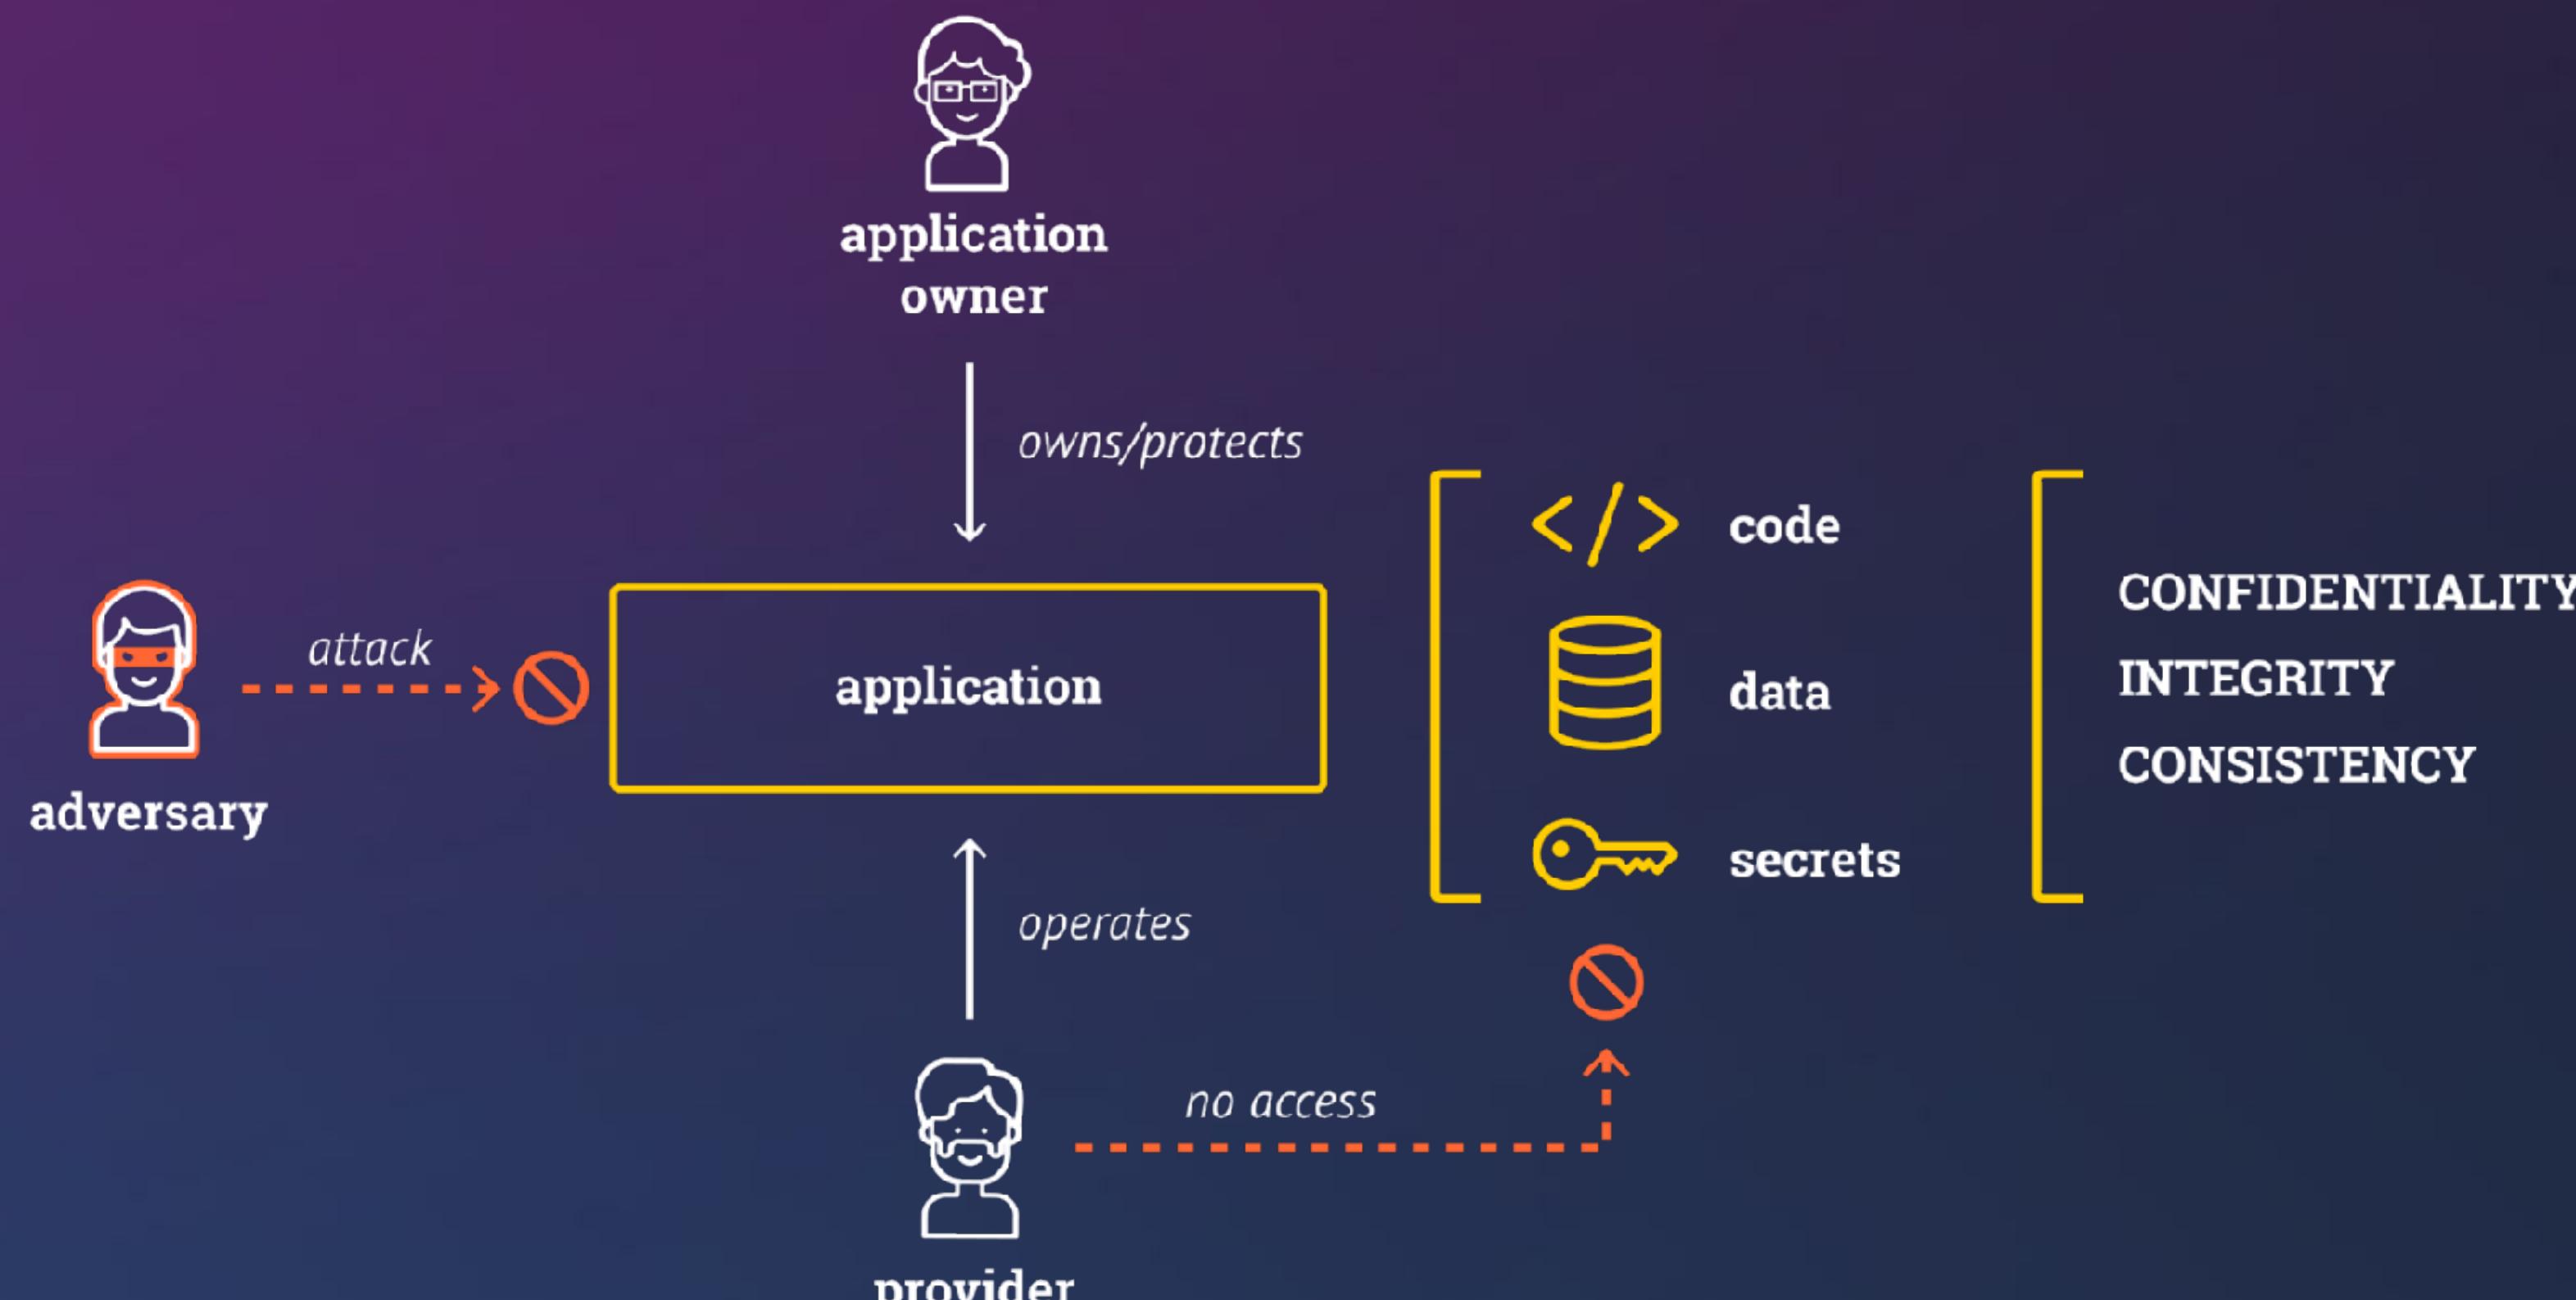
How to define who is authorized
if one cannot trust the operating / hypervisor / ... system?

Protection Goals of Confidential Computing

- **Confidentiality:**
 - only authorized users/programs can read
- **Integrity:**
 - only authorized users/programs can update
- **Consistency**
 - always accessing the last version

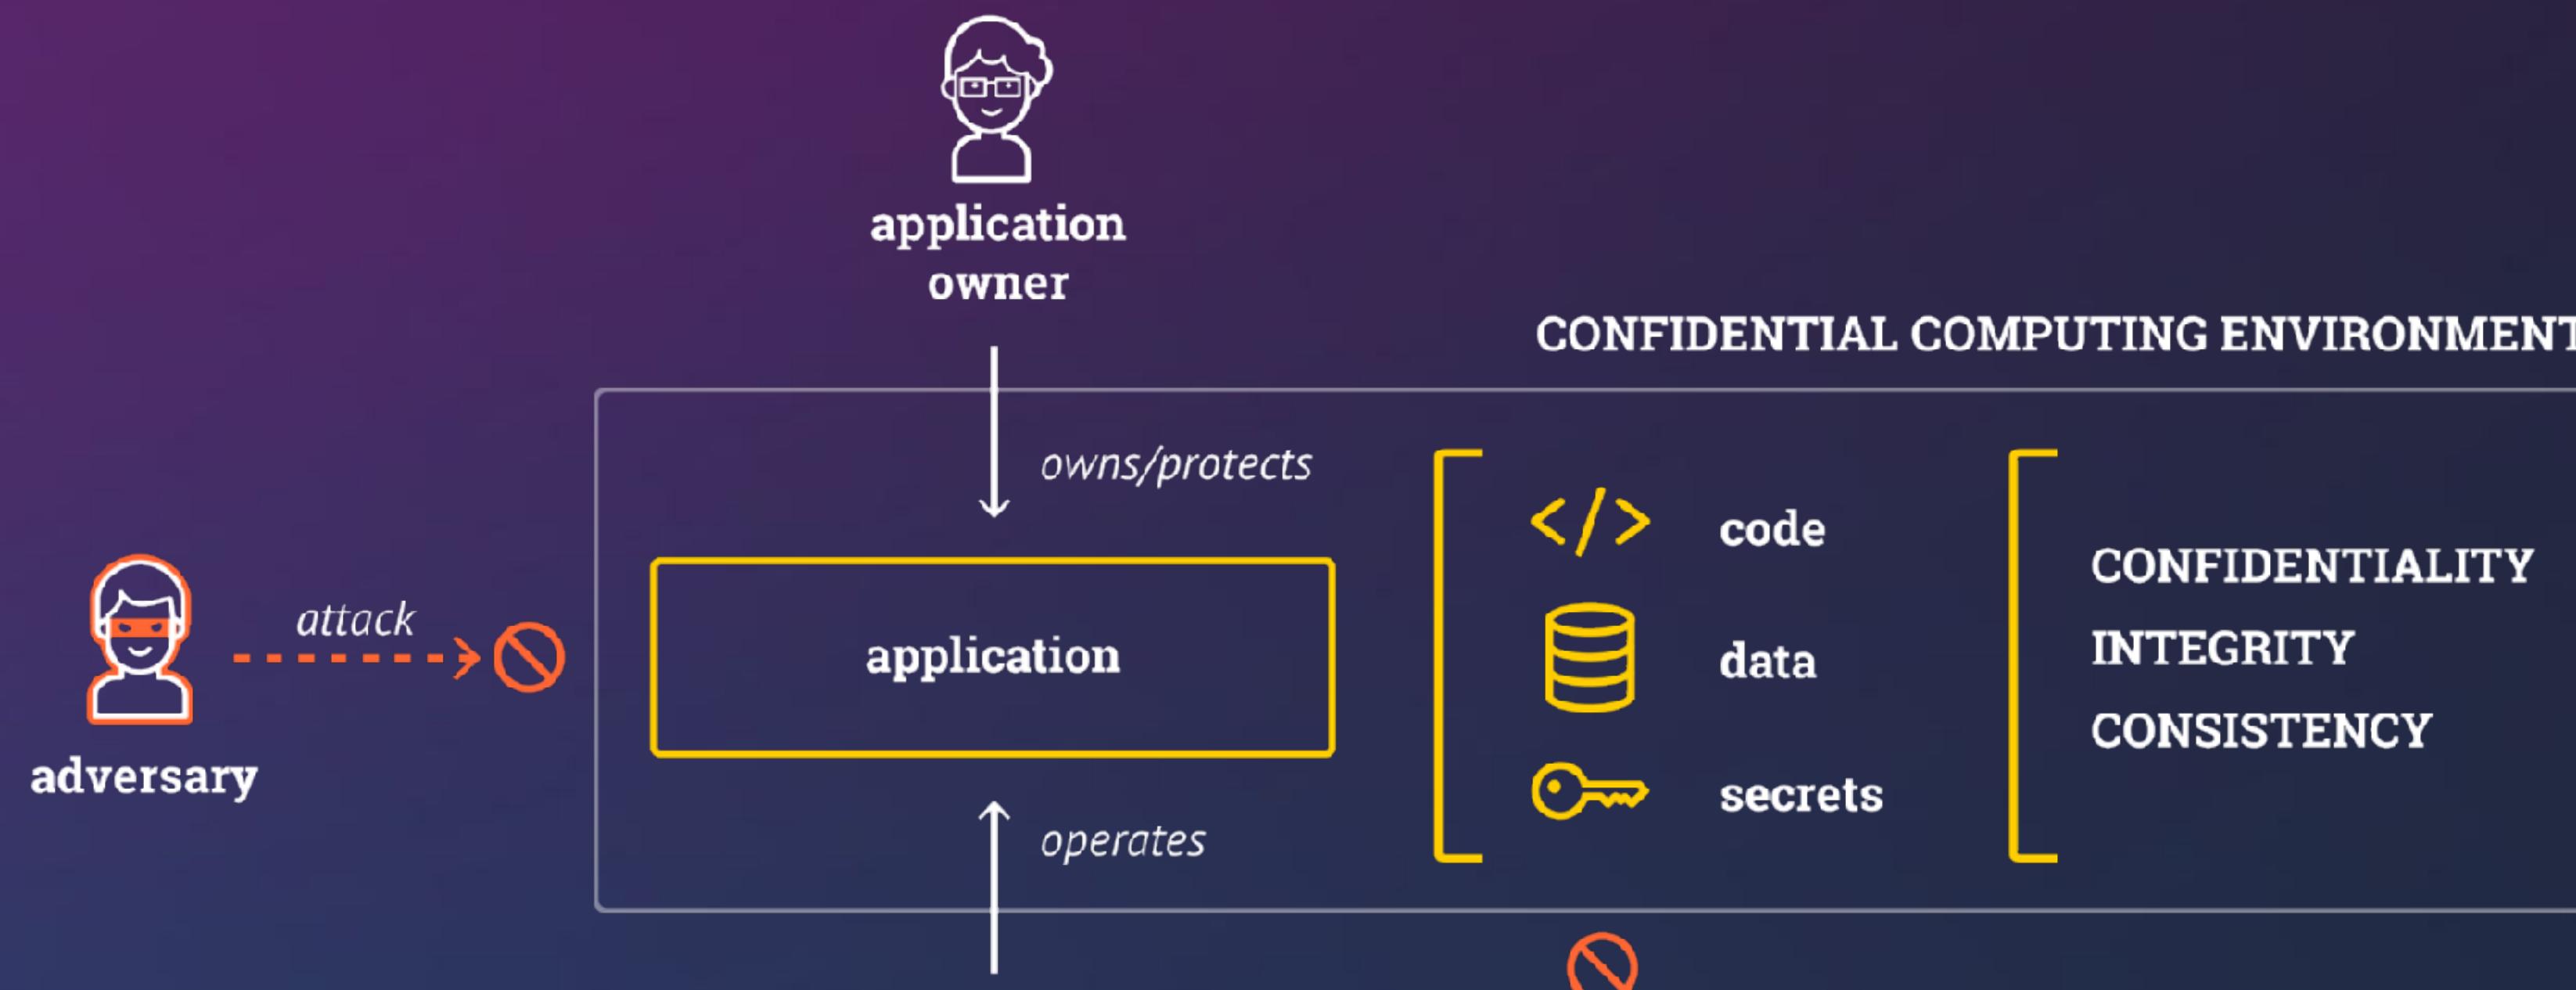


Protect against Adversaries and Cloud Provider!

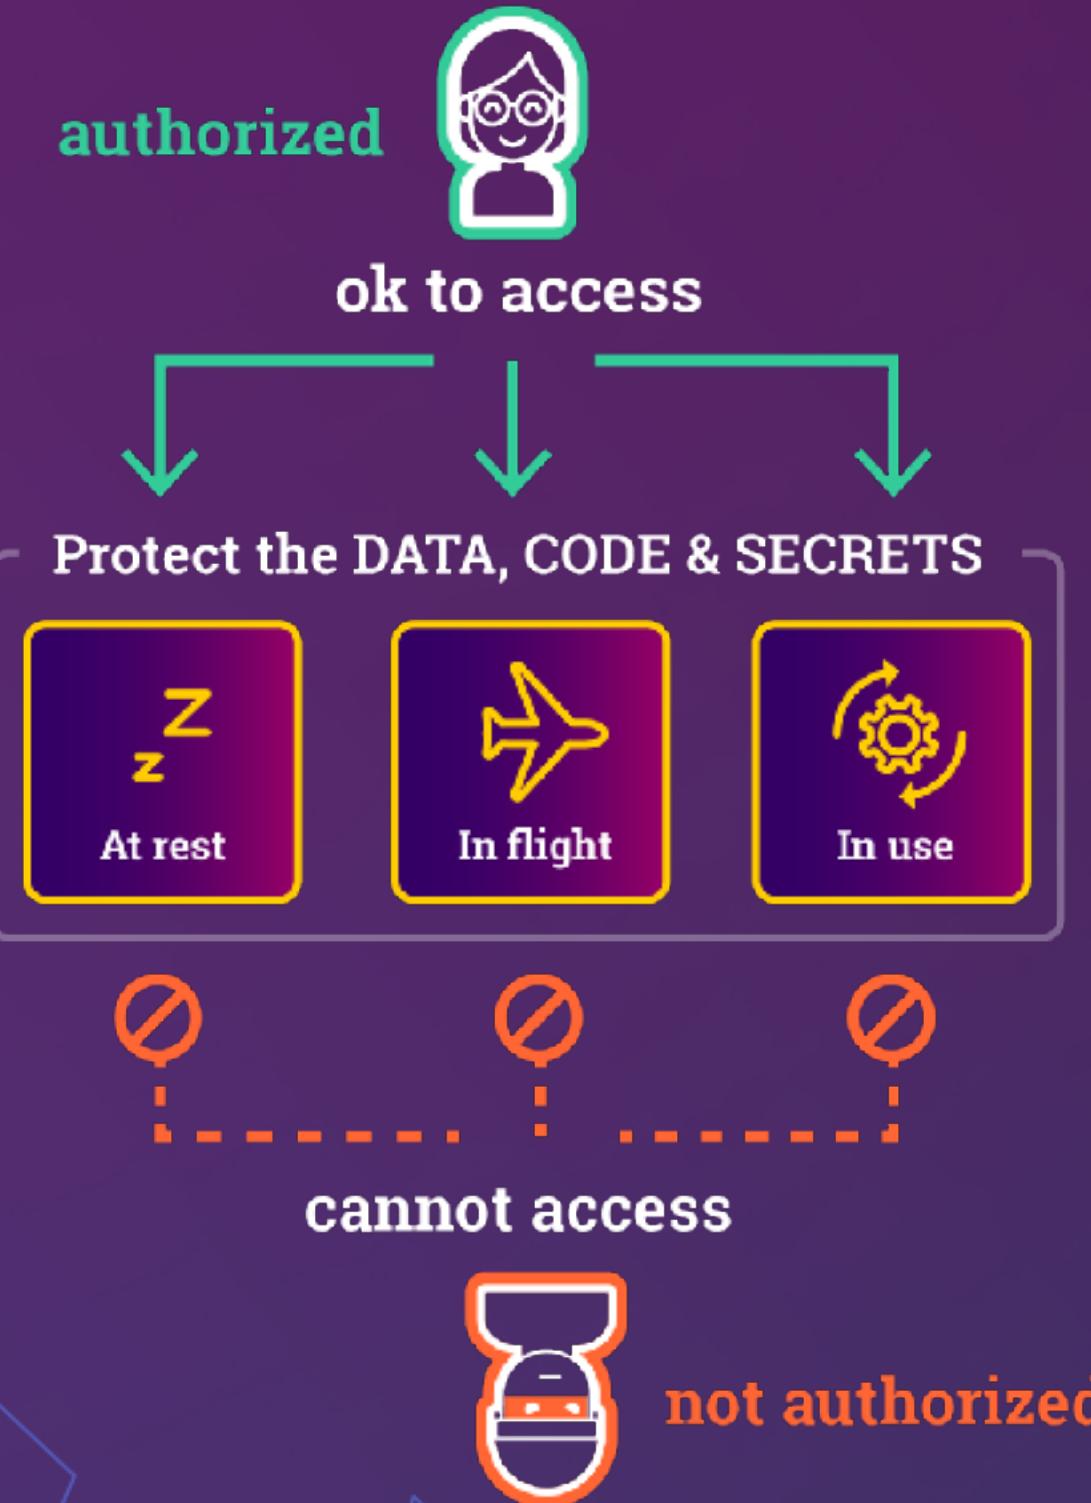


„application-oriented security“

Approach: Confidential Computing



„application-oriented security“



Protection in Use

- i.e., in main memory -

PROBLEM:



How to ensure confidentiality, integrity and consistency if adversary has root and hardware access?

PROBLEM:



How to ensure confidentiality, integrity and consistency if adversary has root and hardware access?

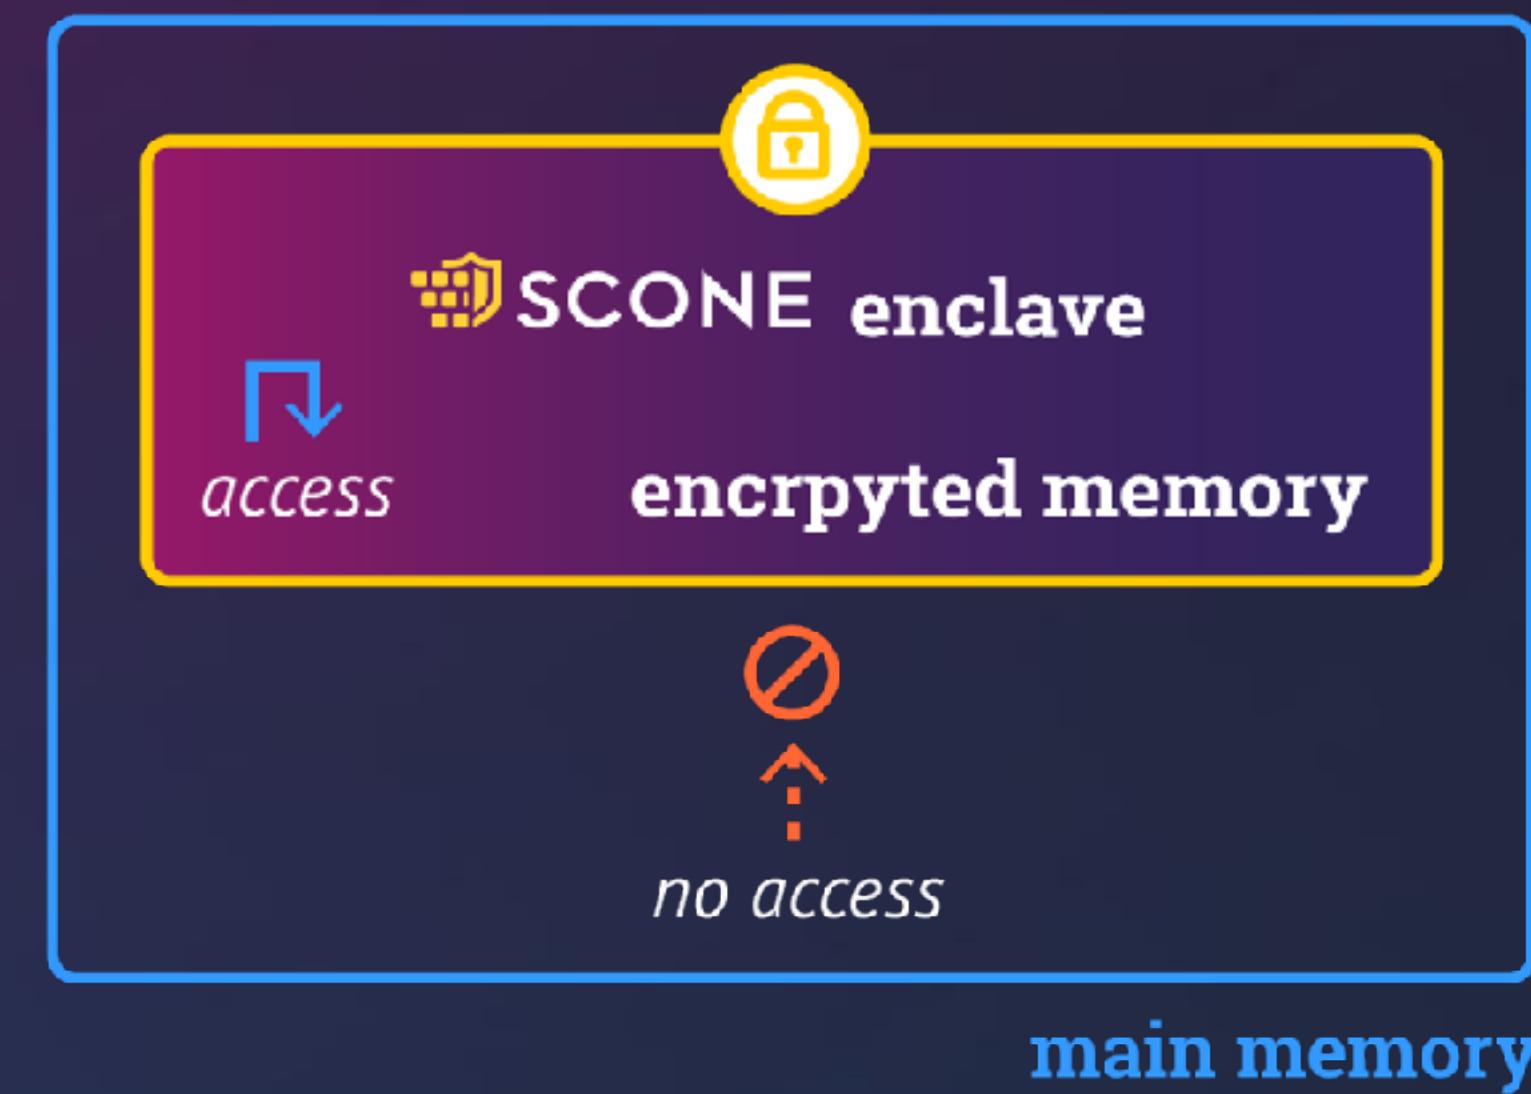
APPROACH:



Run applications in Trusted Execution Environments (TEE)
exclude access by root users!

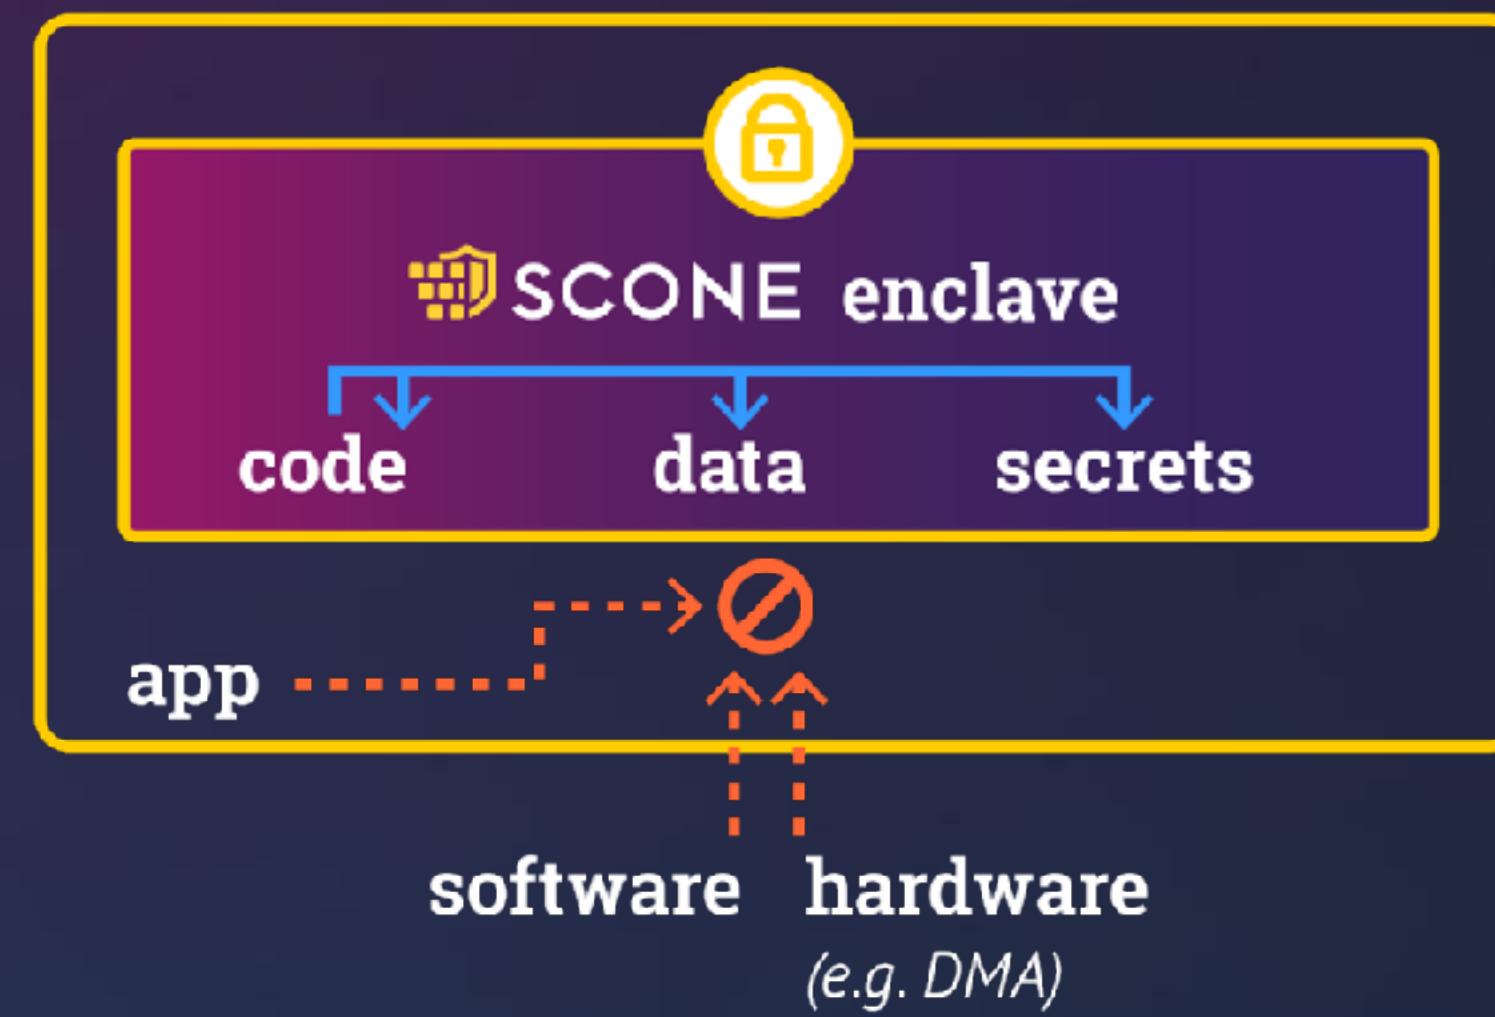
Trusted Execution Environment (TEE)

- Enclaves are a TEE:
 - CPU instruction set extension
 - service runs inside an encrypted memory region, aka, **enclave**
 - **random key** – generated by CPU, or
 - **external key** – given by external entity - **can we trust this key?**
 - only code running inside of enclave can see content
- **Hardware with different granularities:**
 - **enclave**: process runs inside encrypted memory (Intel SGX)
 - **encrypted VM**: a whole VM (AMD SEV, ARM Realms, Intel TDX)
 - can be used to implement enclaves



Enclaves

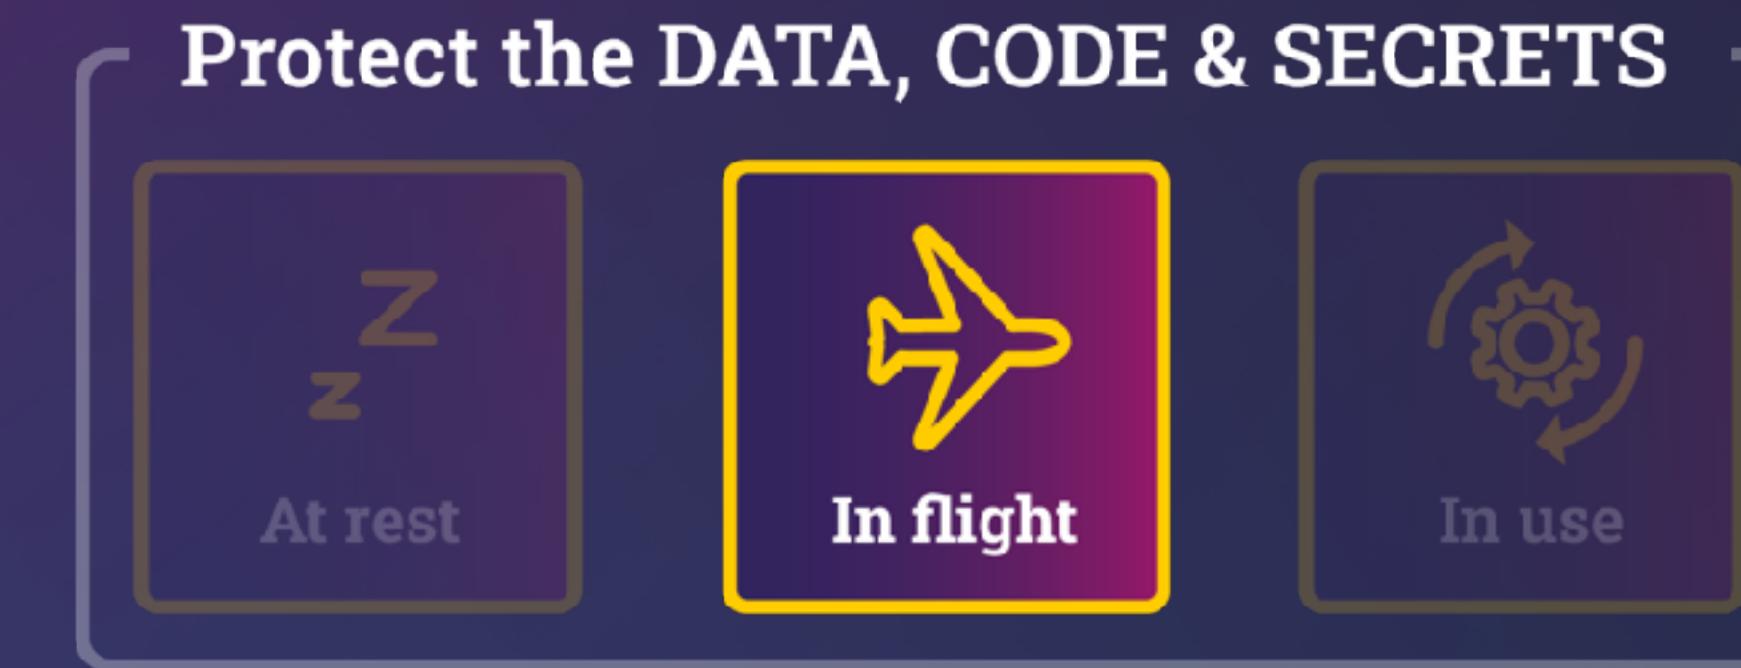
- Protect data/code/secrets in use (i.e, in main memory):
 - run application code in encrypted memory region (aka **enclave**)
 - only code in enclave can access region



CPU extension:
instructions to create enclave



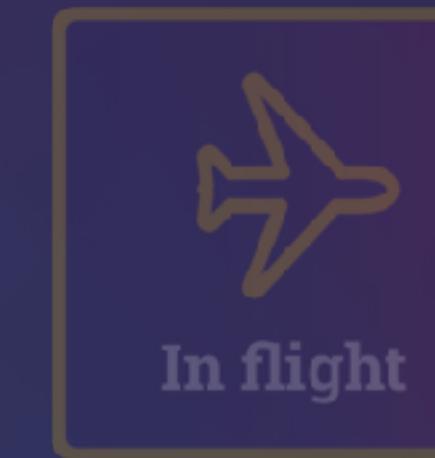
Protecting Data in Flight: SCONE Network Shield



introduced in a later section

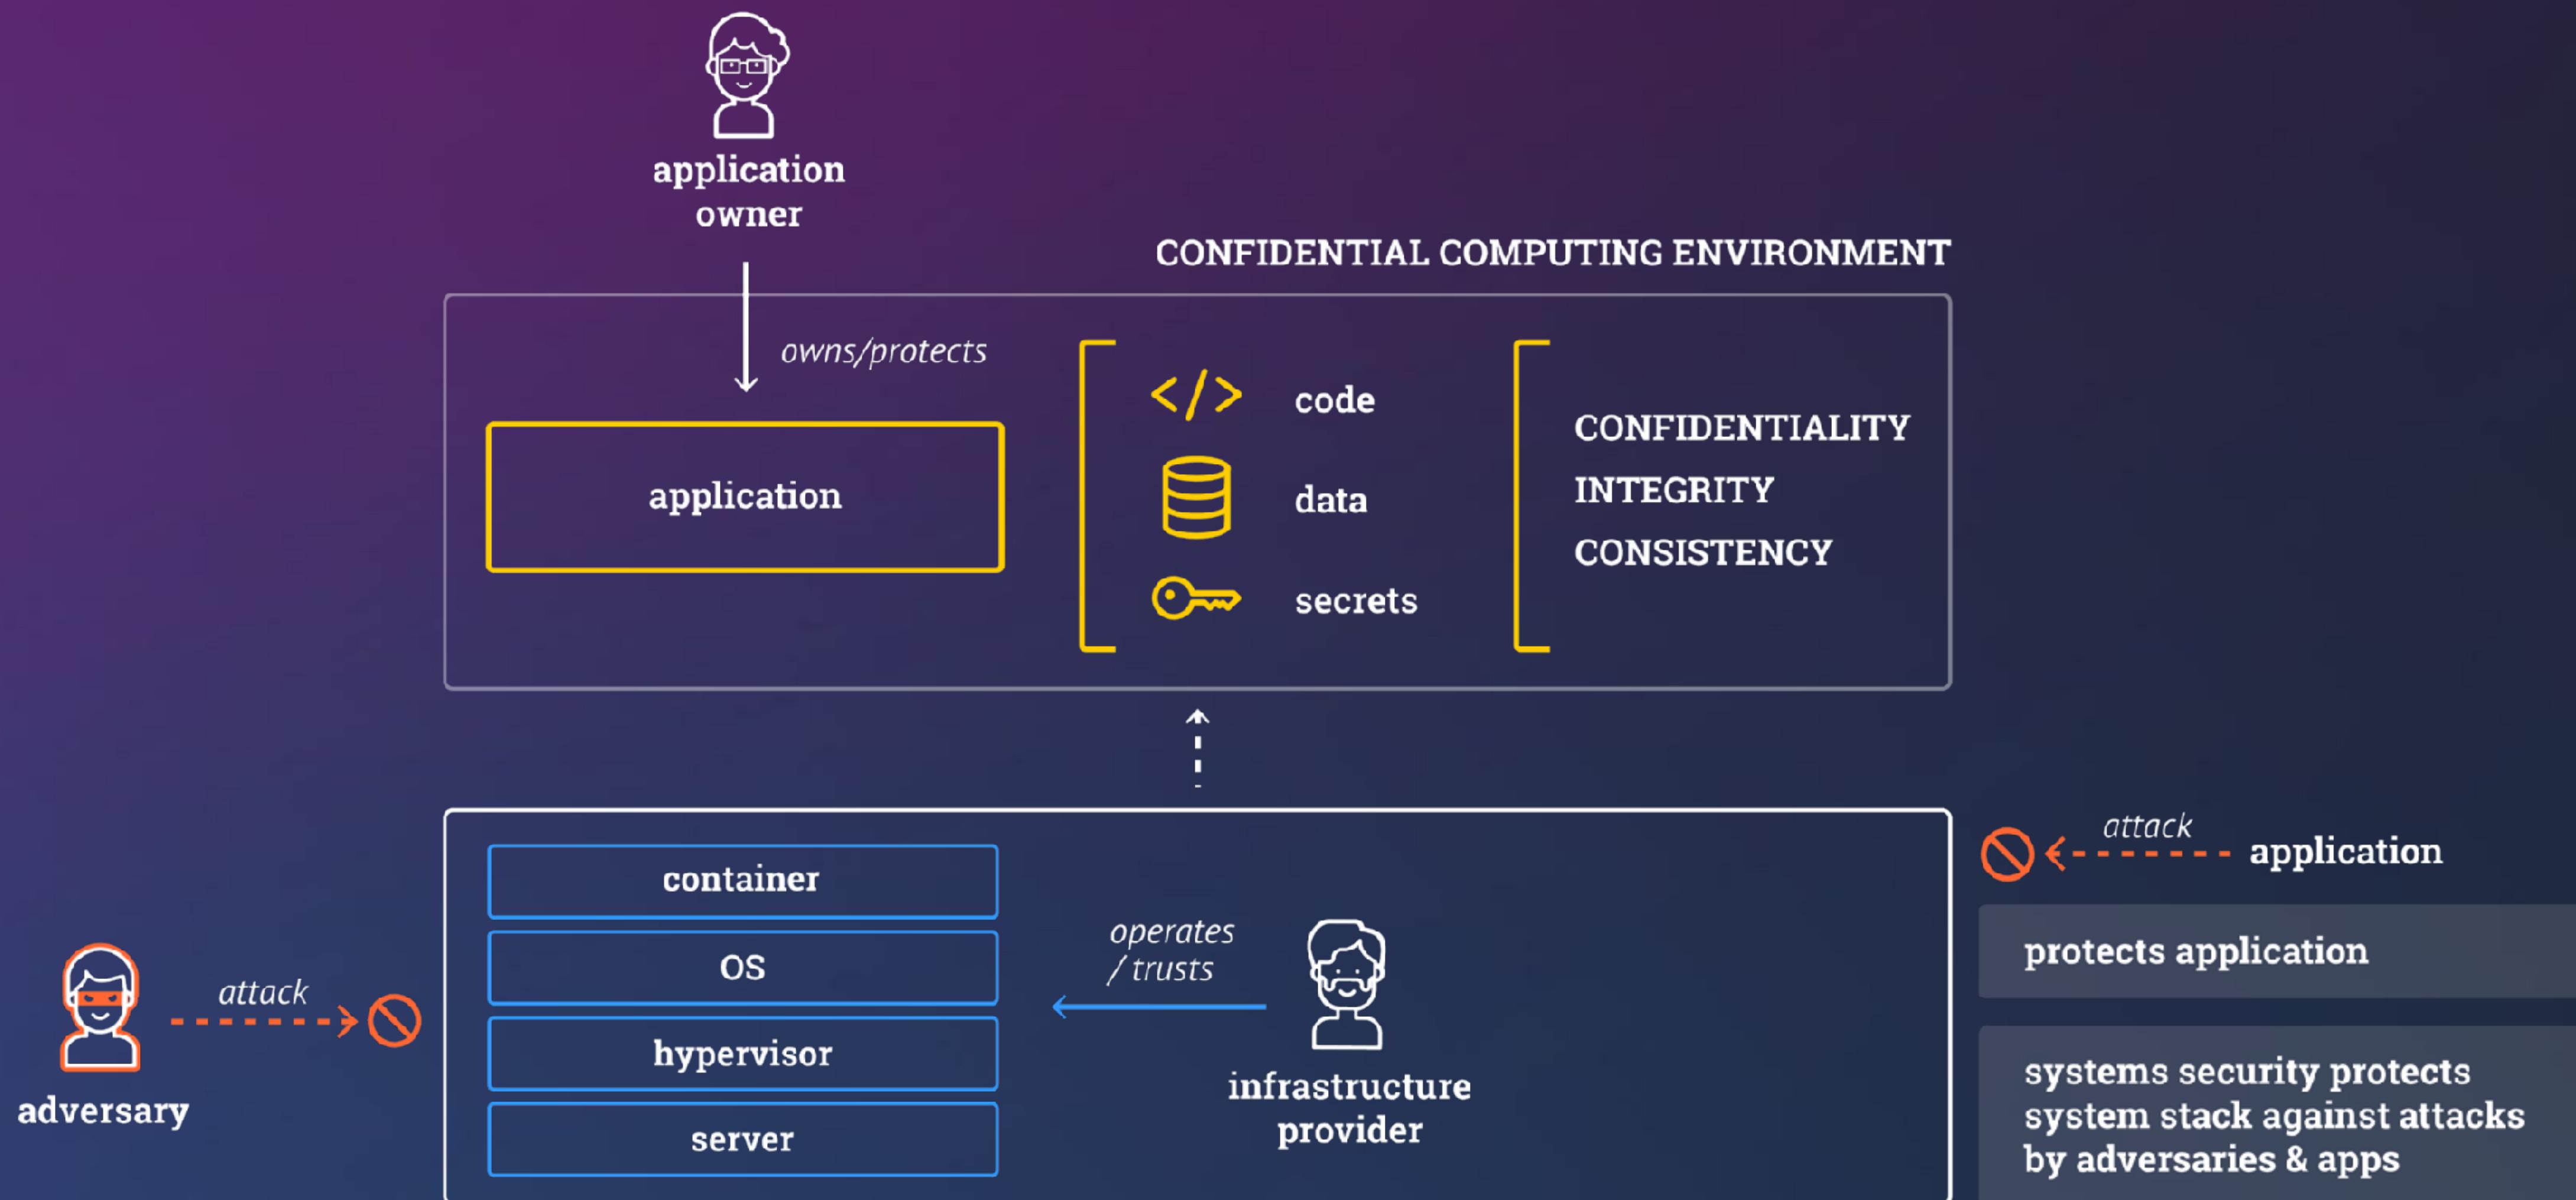
Protecting Data at REST: SCONE Fileshield

Protect the DATA, CODE & SECRETS



introduced in a later section

Complementary: Systems Security + Confidential Computing



Availability

- **DEFINITION:**
 - The probability that a system/ application / service will operate satisfactorily at a given point in time when
- **EXAMPLES:**
 - 99.99% (4 nines availability)
 - 99.999% (5 nines availability)
 - 99.9999% (6 nines availability)

Durability

- **DEFINITION:**
 - The probability that a data item will be accessible after one year
- **EXAMPLES:**
 - 99.99999999% (11 nines durability)

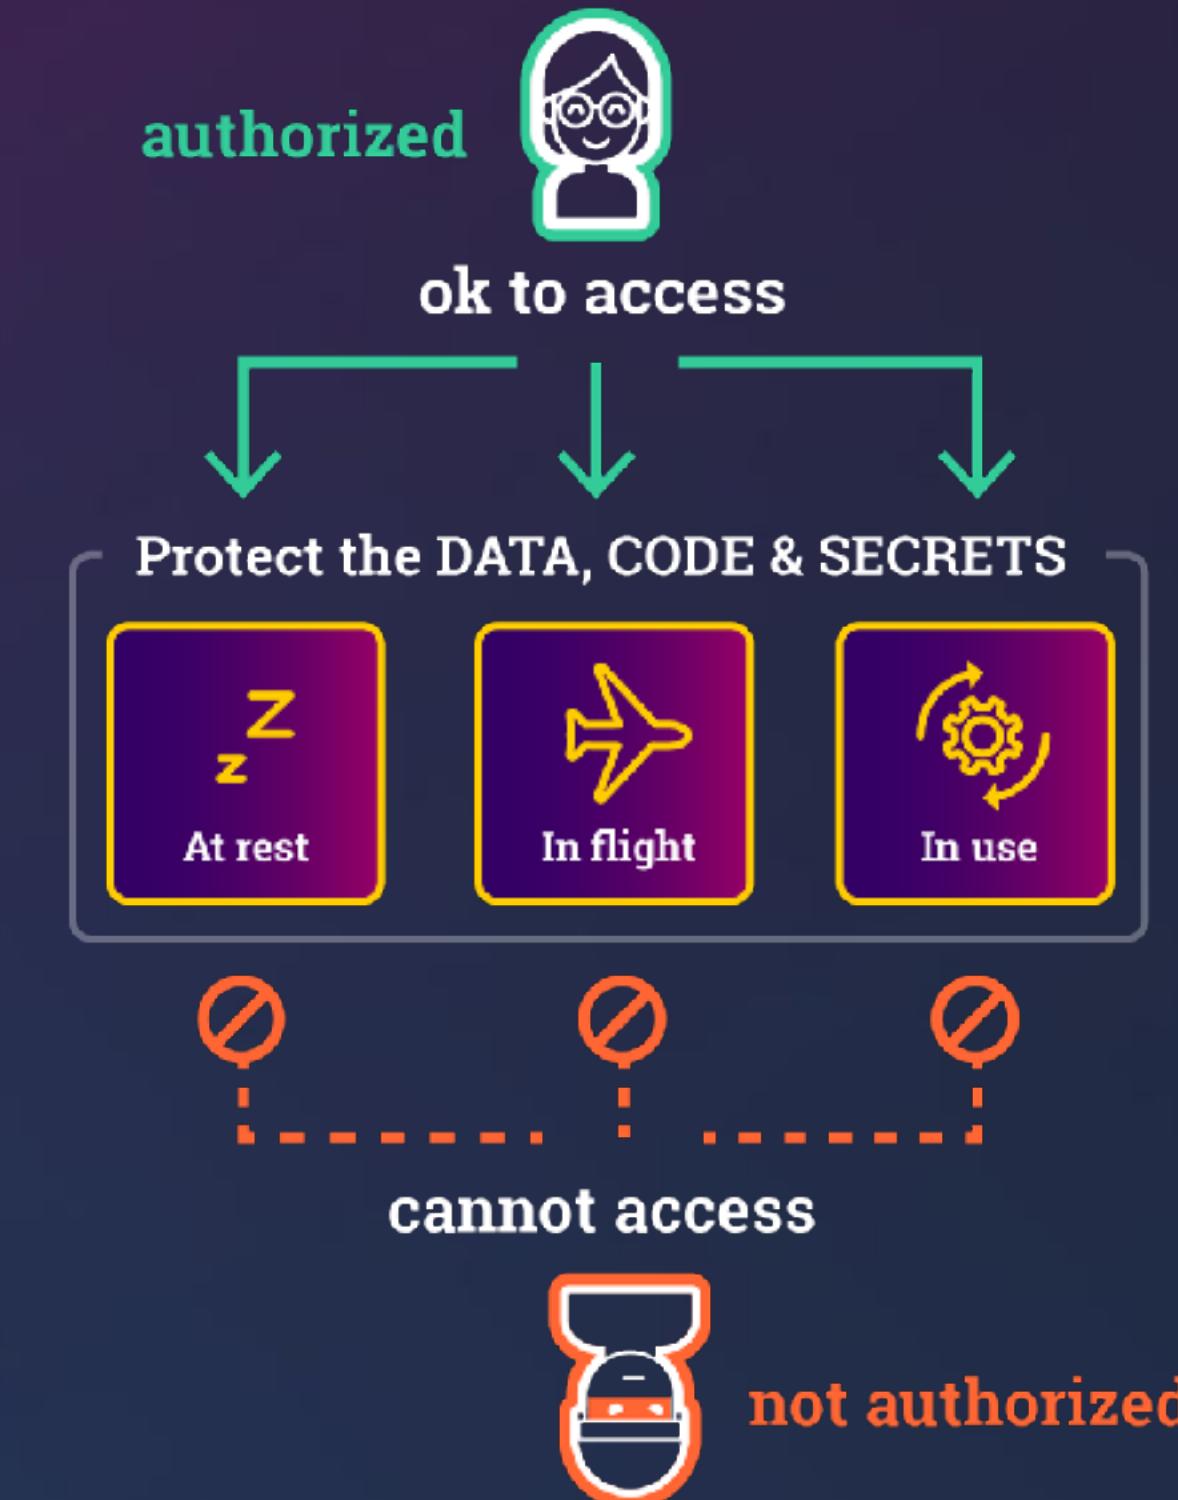
Availability and Durability

- Protection goals NOT addressed by confidential compute:

- **Availability:** the probability that information is available when it is needed
- **Durability:** the probability that information will survive for one year

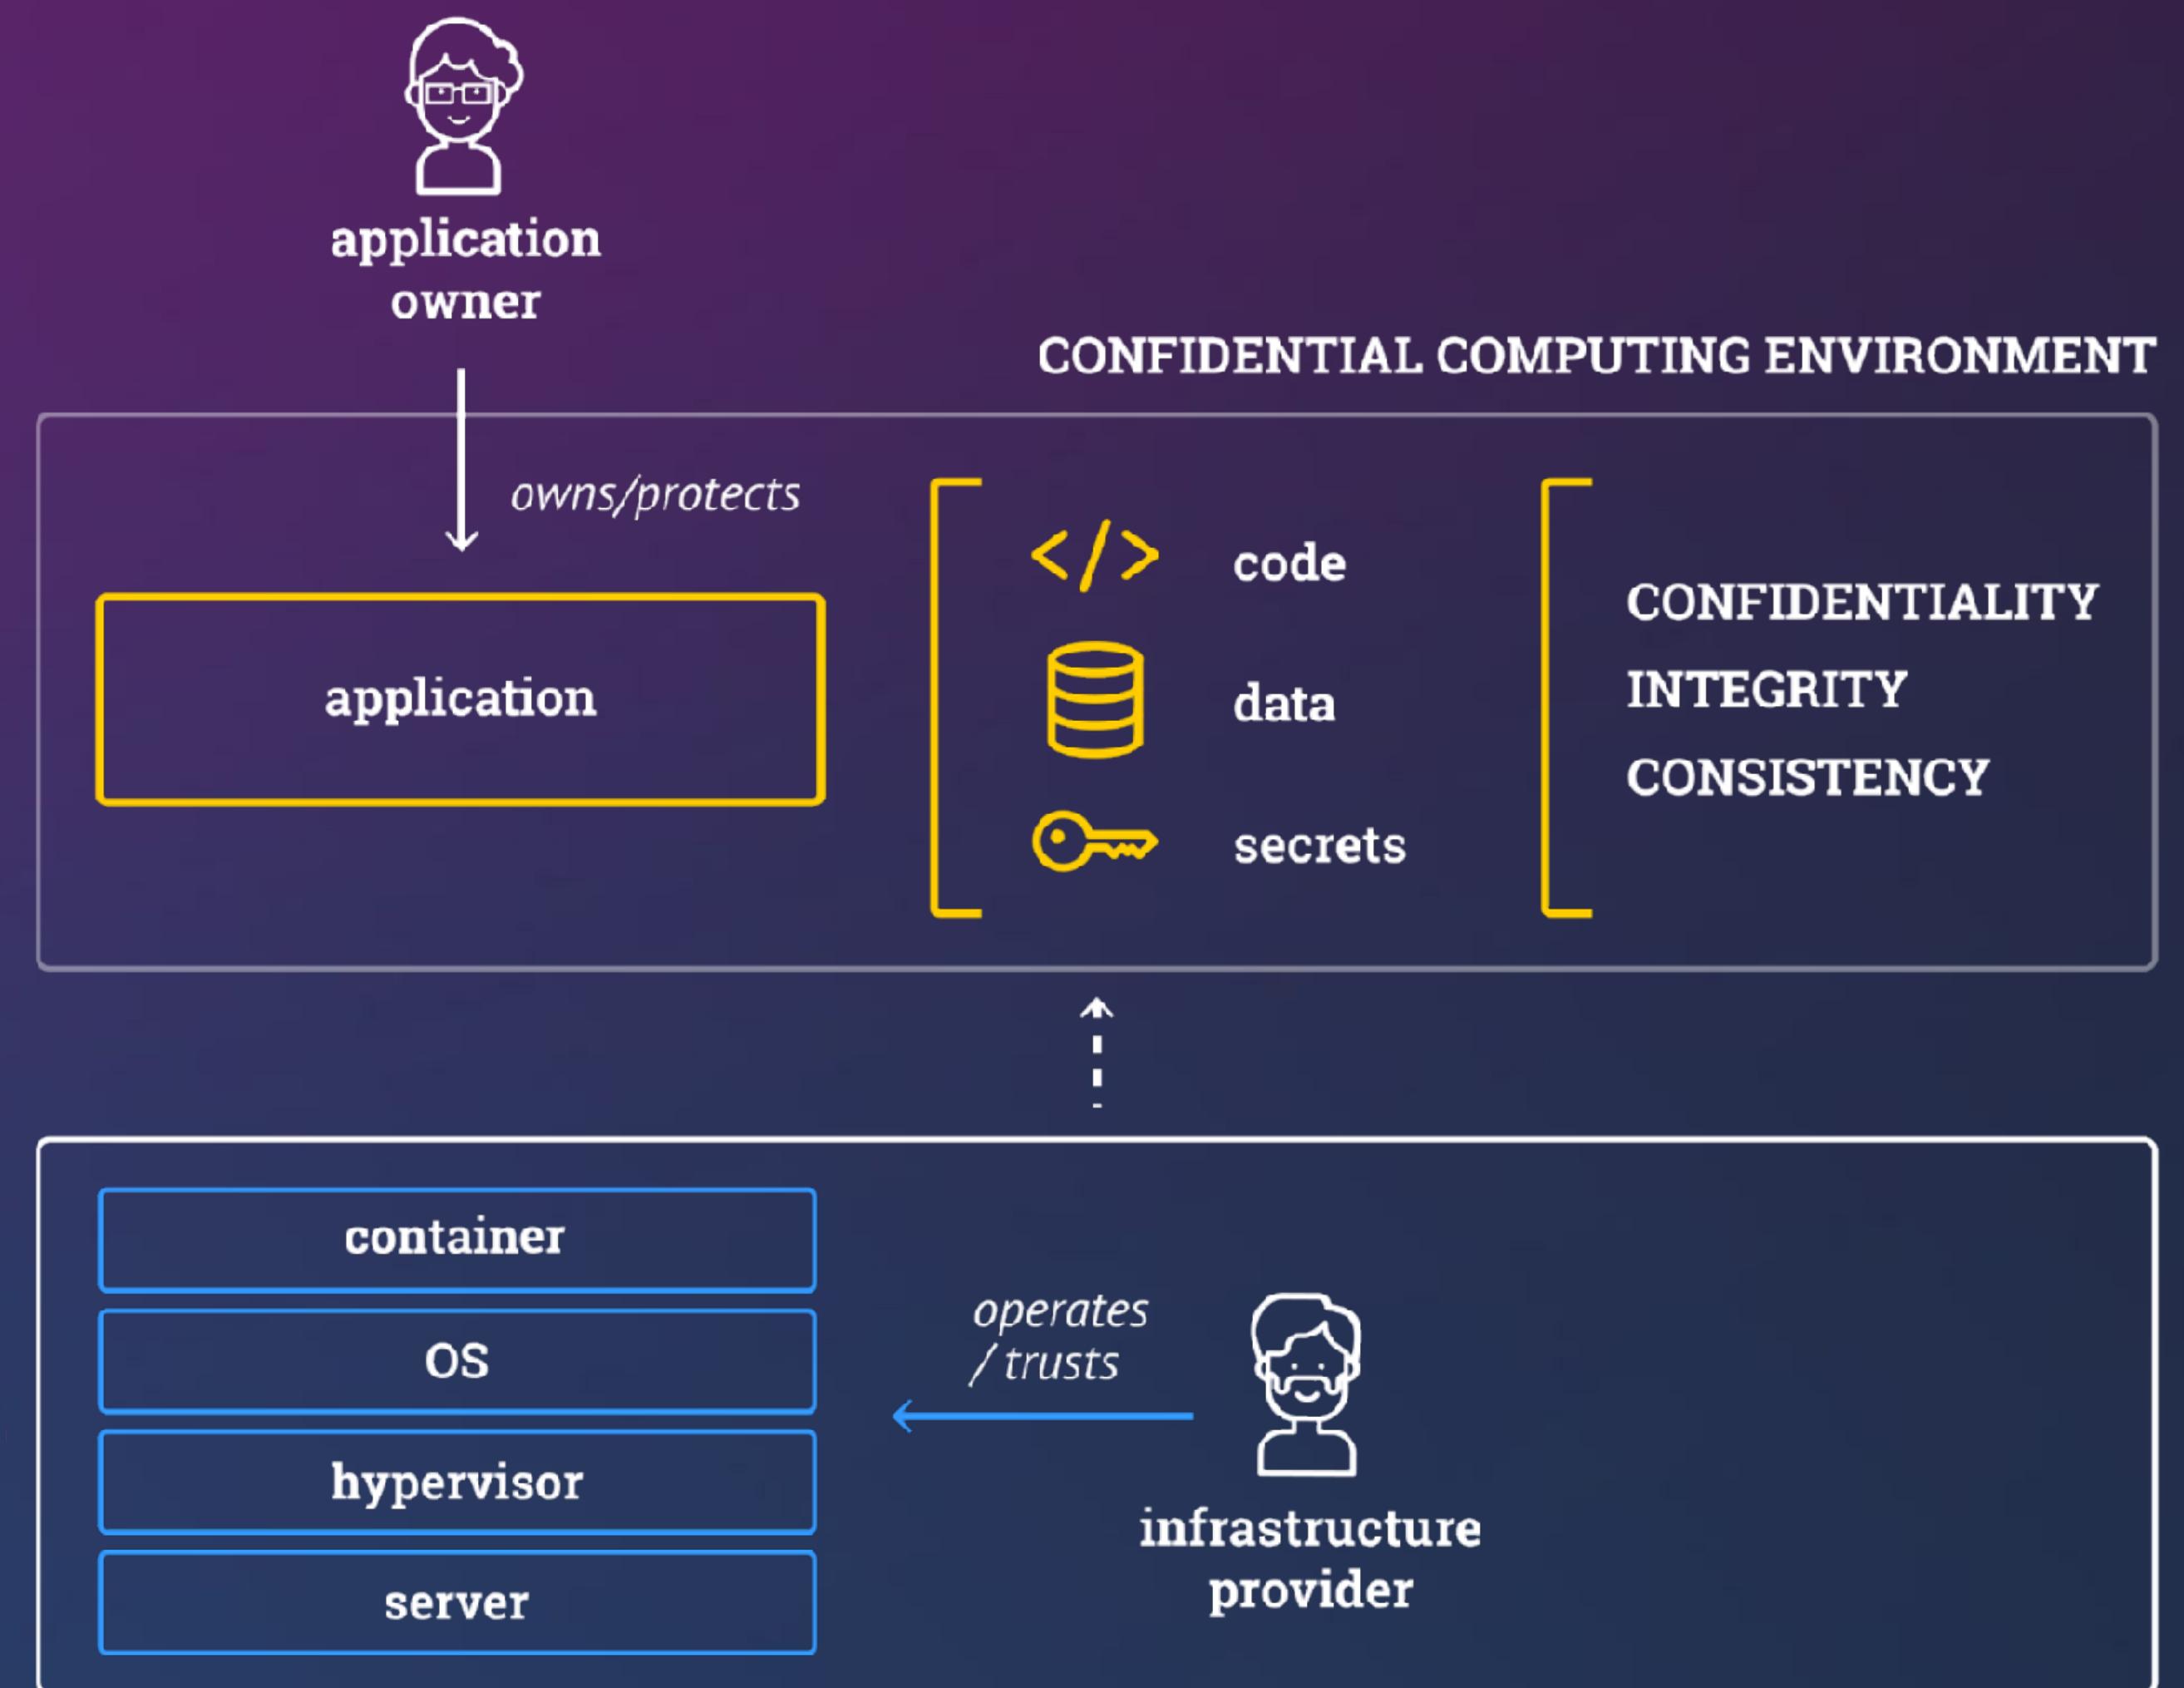
- Approach:

- address this via SLA (Service Level Agreements) of provider
- use **automatic reconciliation** of the services



introduced in a later section

Availability



General Approach:

we measure **availability** and **durability** & define SLA with provider

restart by Kubernetes! 

Implementation:

we use reconciliation to improve availability of services

Note:

Confidential Computing does not address availability nor durability

(SLA = Service Level Agreement)

Example



What information to protect?

- Protection of
 - **Code**, e.g., modern AI programs written in Python
 - **Data**, e.g., training data to create AI models
 - **Secrets**, e.g., key used to encrypt database

What information to protect?

- Protection of
 - **Code**, e.g., modern AI programs written in Python
 - **Data**, e.g., training data to create AI models
 - **Secrets**, e.g., key used to encrypt database
- **Example:**
 - *MariaDB* supports encryption of database
 - encryption key is stored in configuration file
 - configuration file protected via access control:
 - i.e., can be read and written by MariaDB (user) as well as any root (=privileged) user

Is Operating System access control sufficient?

- Sufficient if we define that
 - only **authorized user** can become root users

Is access control sufficient?

- Sufficient if we define that
 - only **authorized user** can become root users
- What about
 - **adversary** gaining root access (e.g., stealing credentials)?
 - **authorized user** laid off -> could become an **adversary**?

Is access control sufficient?

- Sufficient if we define that
 - only **authorized user** can become root users
- What about
 - **adversary** gaining root access (e.g., stealing credentials)?
 - **authorized user** laid off -> could become an **adversary**?
- **Approach:**
 - Define **threat model** that defines the power of the adversary